Demystifying ChatGPT

A shallow deep-dive into LLMs

SKYLYZE DSA

DSA Daten- und Systemtechnik GmbH | SKYLYZE | Roland Stoffel | Pascal Rößner

AGENDA

- Introduction ChatGPT aka LLM (Lecture: 15min)
  - Defintion and Use-Cases
  - Theoretical Background
  - Project Lifecycle
- From Scratch (Hands-On Session: 35-45min)
  - Train your own model
- Challenges of LLMS and Generative AI (Lecture: 10min)
  - Infrastructure and Costs
  - Transfer to Image Generation
- Experimentation (Hands-On Session: 20min)
  - Have Fun with Stable Diffusion

WHAT IS A LLM

- ChatGPT is a so-called Large Language Model (LLM)
- Large Language Models are deep learning models trained on huge datasets to perform NLP tasks.
- Its core objective is to learn and understand human languages precisely.
- Large Language Models enable the machines to interpret languages just like the way we, as humans, interpret them.

Name	Release Date	Parameter Size
GPT-4	2023	4.6 trillion
LLaMA	2022	1.5 trillion
FLAN UL2	2022	1.3 trillion
BLOOM	2022	176 billion
LaMDA	2021	173 billion
MT-NLG	2020	530 billion
GPT-3	2020	175 billion
GPT-2	2019	1.5 billion
BERT	2018	340 million
ELMo	2017	94 million

HISTORICAL EVOLUTION

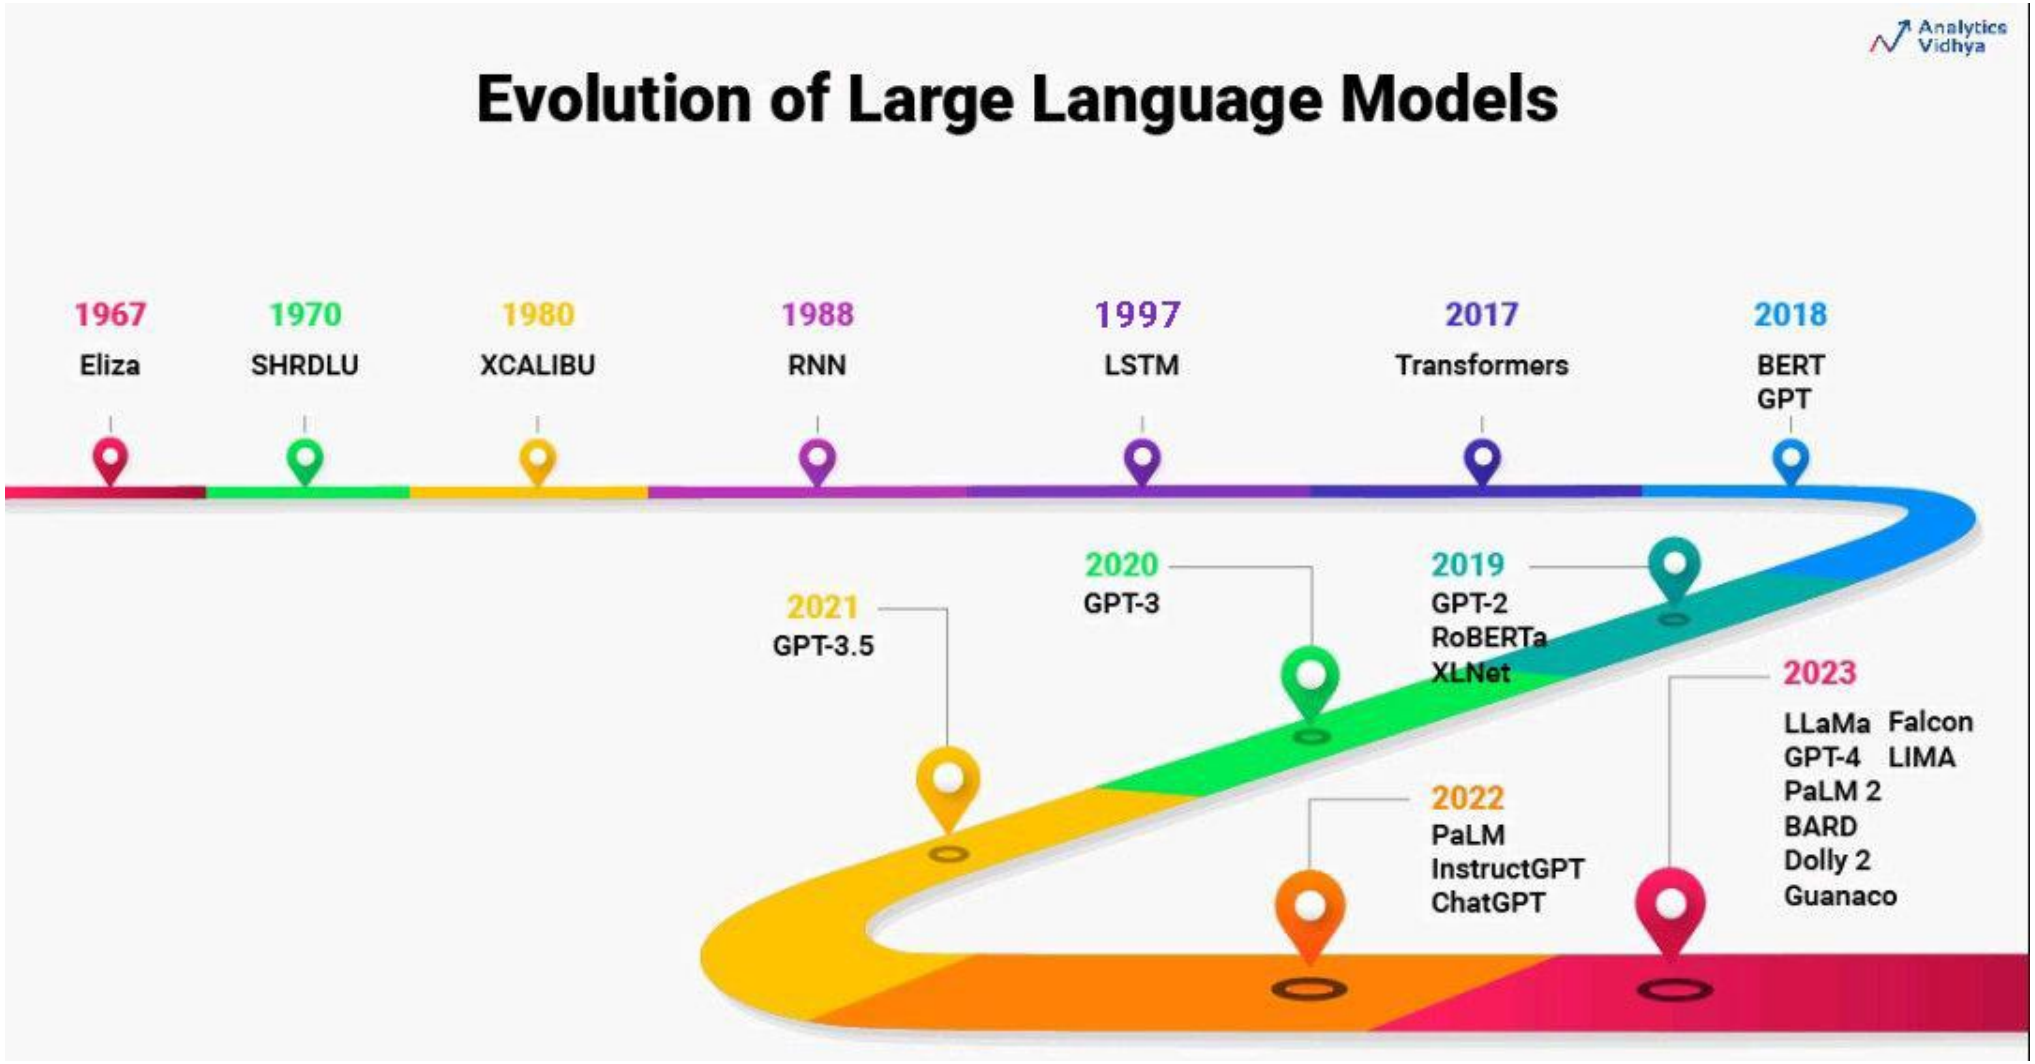

PROMPTS AND COMPLETIONS

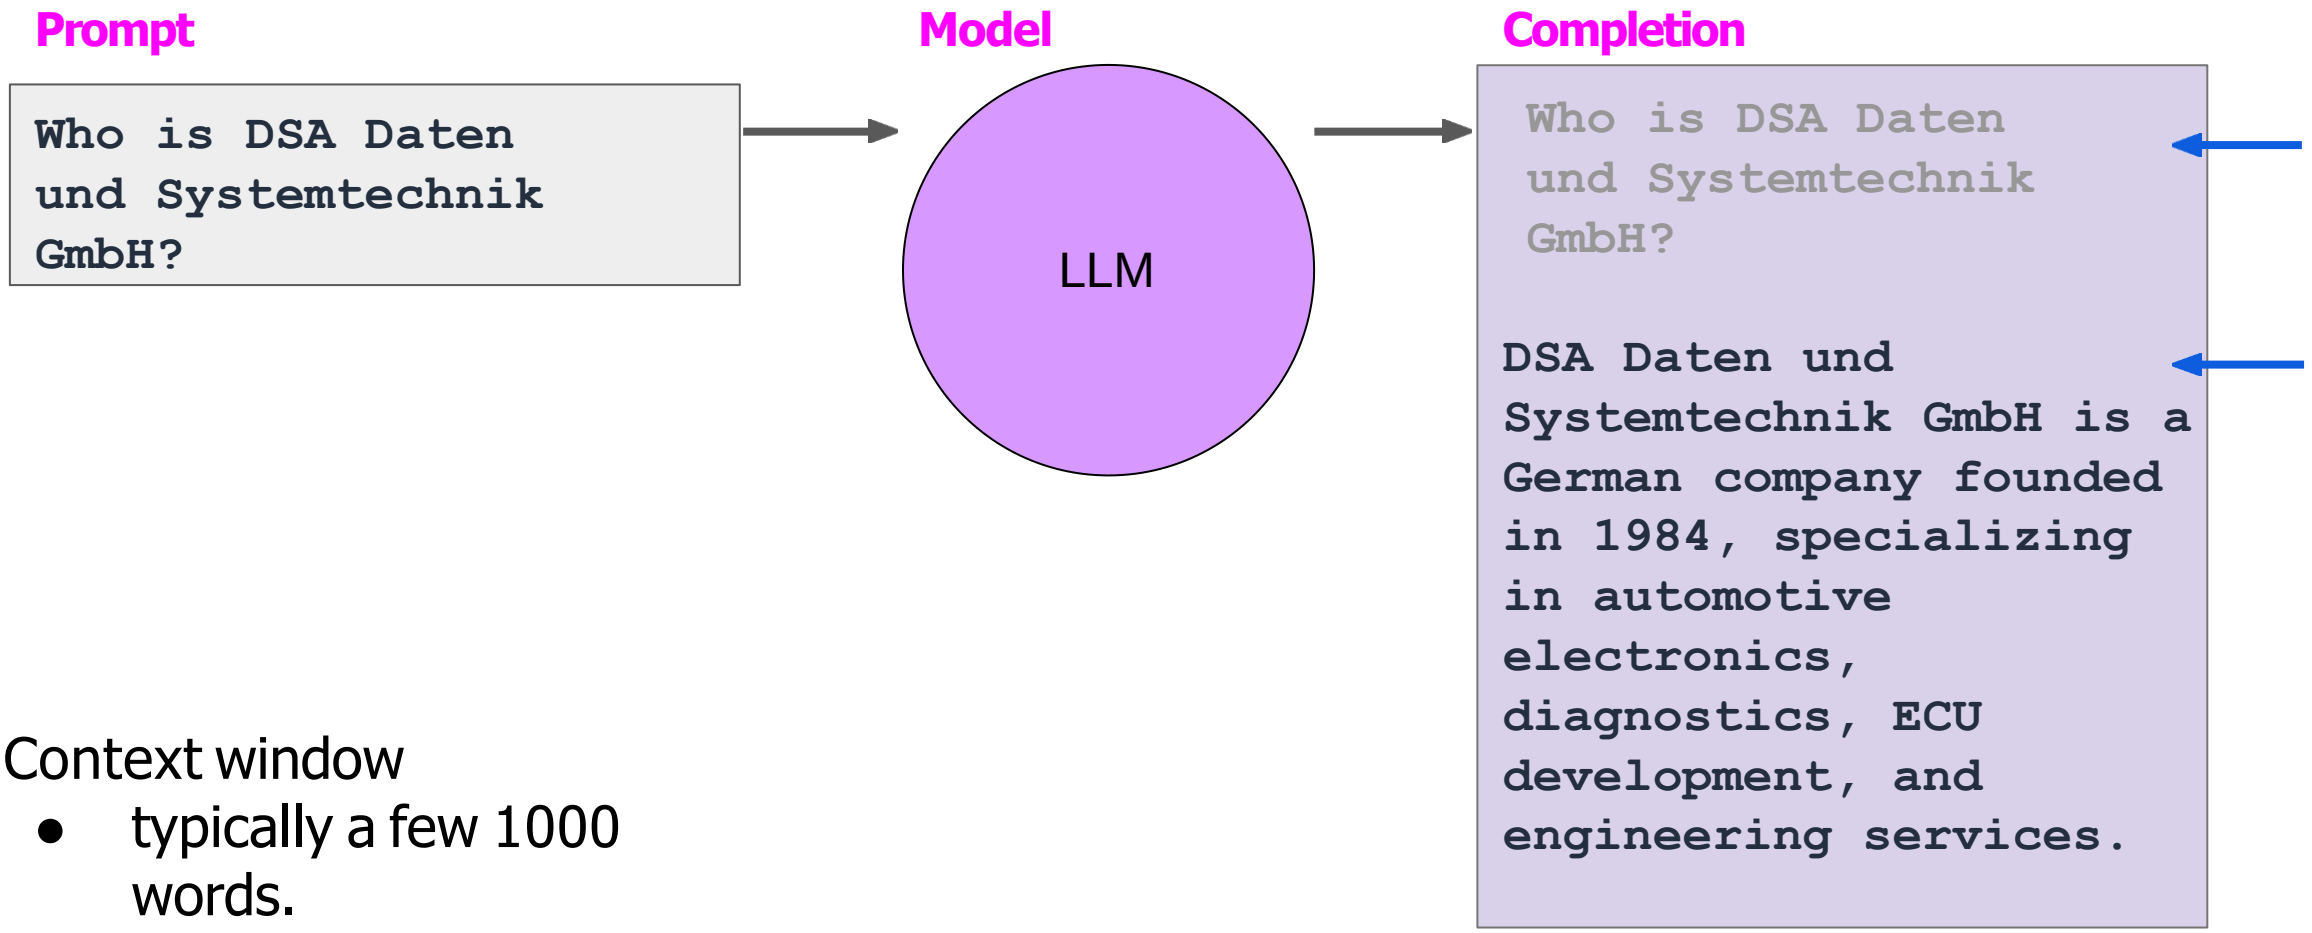

Context window

- typically a few 1000 words.

UNDERSTANDING LANGUAGE CAN BE CHALLENGING

The teacher's book?

The teacher taught the student with the book.

The student's book?

TRANSFORMER ARCHITECTURE

Attention Is All You Need

Ashish Vaswani* Google Brain avaswani@google.com

Noam Shazeer* Google Brain noam@google.com

Niki Parmar* Google Research nikip@google.com

Jakob Uszkoreit* Google Research usz@google.com

Llion Jones* Google Research llion@google.com

Aidan N. Gomez* † University of Toronto aidan@cs.toronto.edu

Łukasz Kaiser* Google Brain lukaszkaiser@google.com

Illia Polosukhin* ‡ illia.polosukhin@gmail.com

Abstract

The dominant sequence transduction models are based on complex recurrent or convolutional neural networks that include an encoder and a decoder. The best performing models also connect the encoder and decoder through an attention mechanism. We propose a new simple network architecture, the Transformer, based solely on attention mechanisms, dispensing with recurrence and convolutions entirely. Experiments on two machine translation tasks show these models to

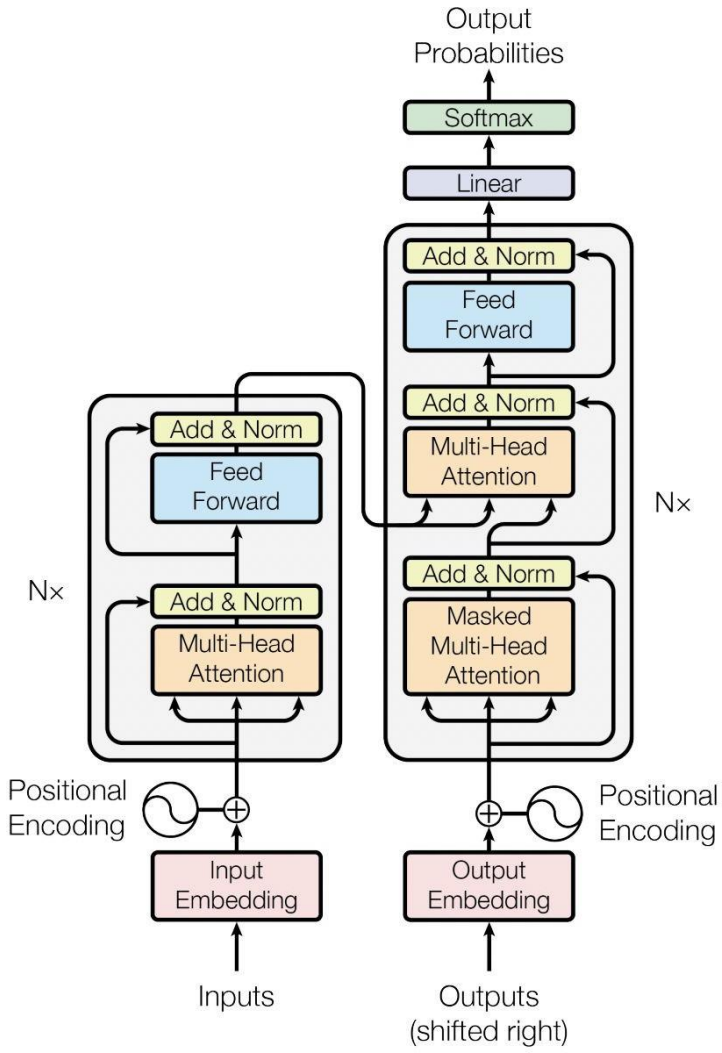

TRANSFORMER ARCHITECTURE

Attention Is All You Need

Ashish Vaswani* Google Brain avaswani@google.com

Noam Shazeer* Google Brain noam@google.com

Niki Parmar* Google Research nikip@google.com

Jakob Uszkoreit* Google Research usz@google.com

Llion Jones* Google Research llion@google.com

Aidan N. Gomez* † University of Toronto aidan@cs.toronto.edu

Łukasz Kaiser* Google Brain lukaszkaiser@google.com

Illia Polosukhin* ‡ illia.polosukhin@gmail.com

Abstract

The dominant sequence transduction models are based on complex recurrent or convolutional neural networks that include an encoder and a decoder. The best performing models also connect the encoder and decoder through an attention mechanism. We propose a new simple network architecture, the Transformer, based solely on attention mechanisms, dispensing with recurrence and convolutions entirely. Experiments on two machine translation tasks show these models to

- Scale efficiently
- Parallelprocess
- Attention to input meaning

ATTENTION IS ALL YOU NEED

The teacher taught the student with the book.

ATTENTION IS ALL YOU NEED

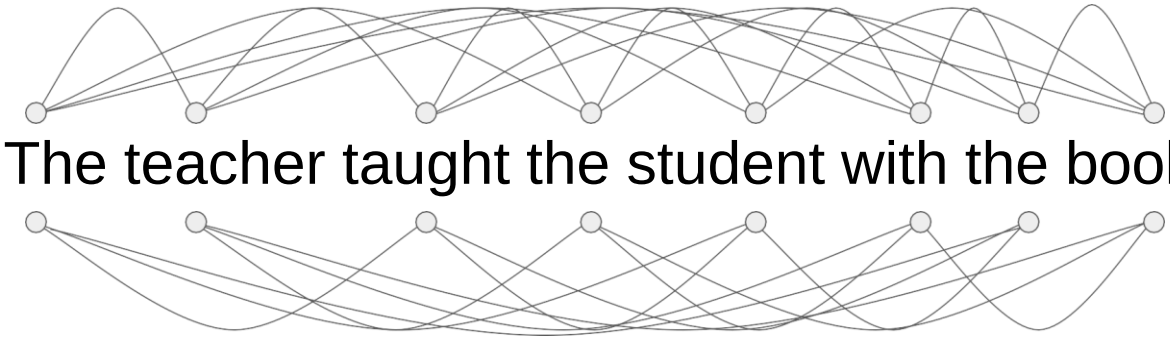

ATTENTION IS ALL YOU NEED

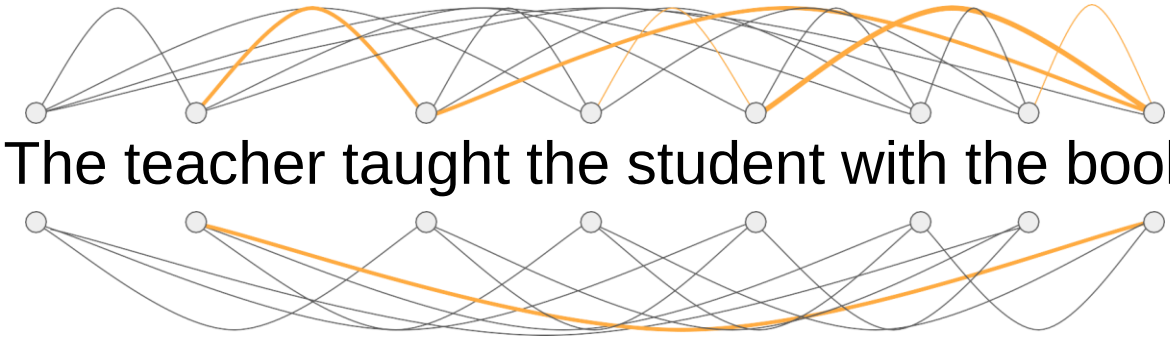

SELF-ATTENTION

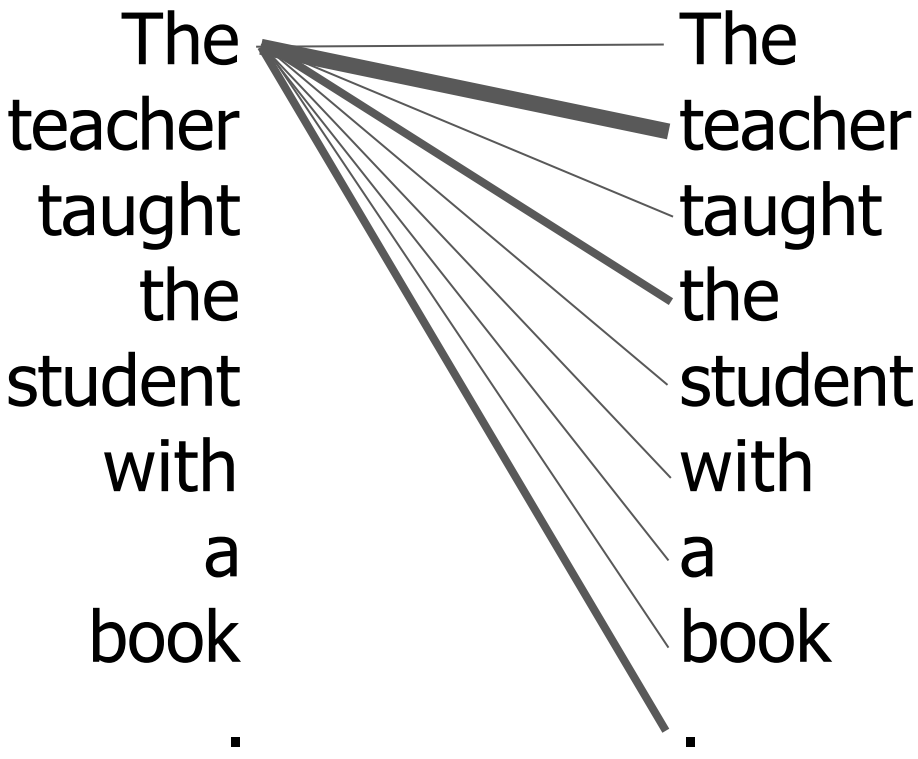

SELF-ATTENTION

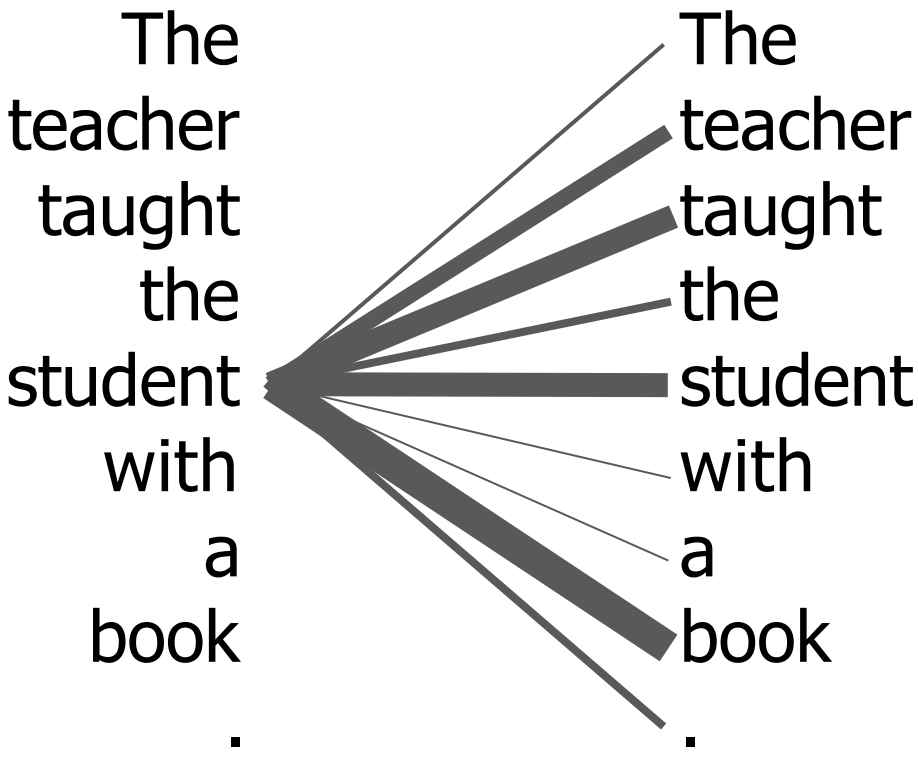

SELF-ATTENTION

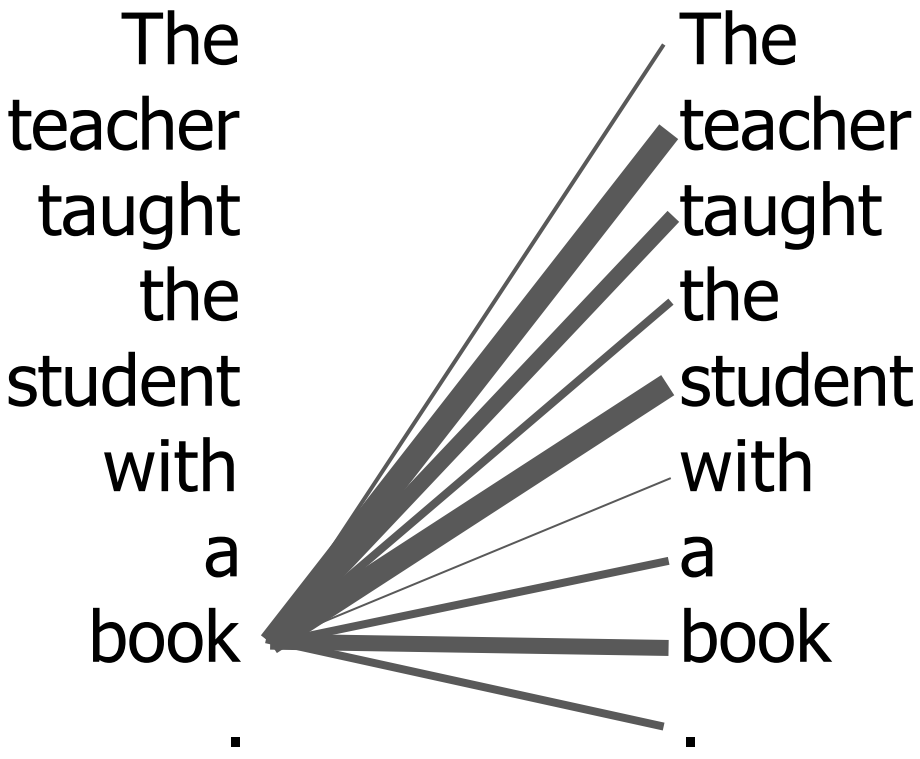

SELF-ATTENTION

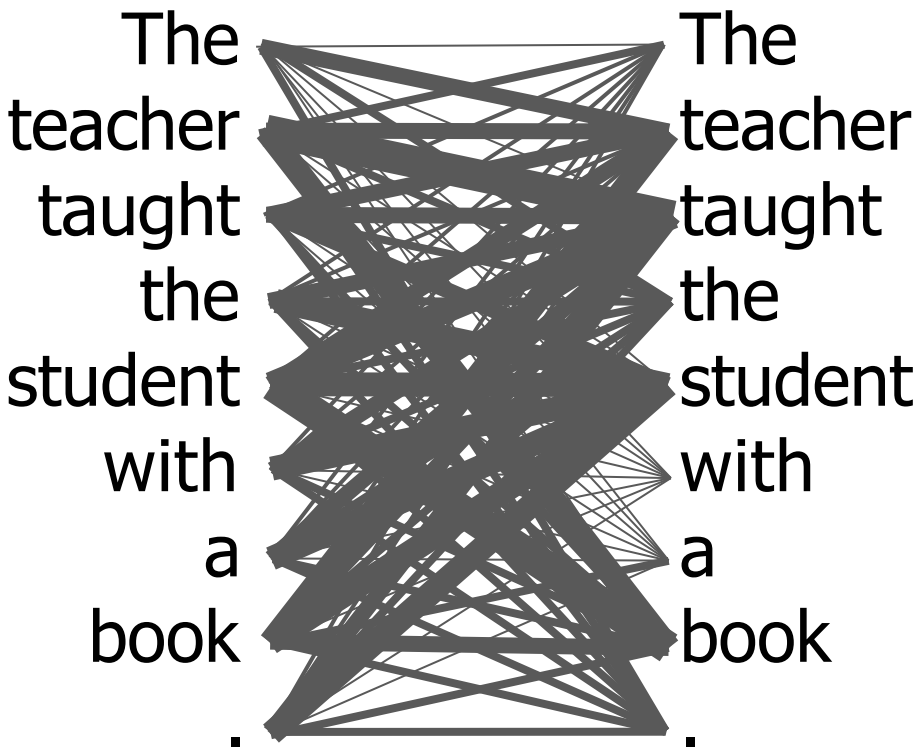

SELF-ATTENTION

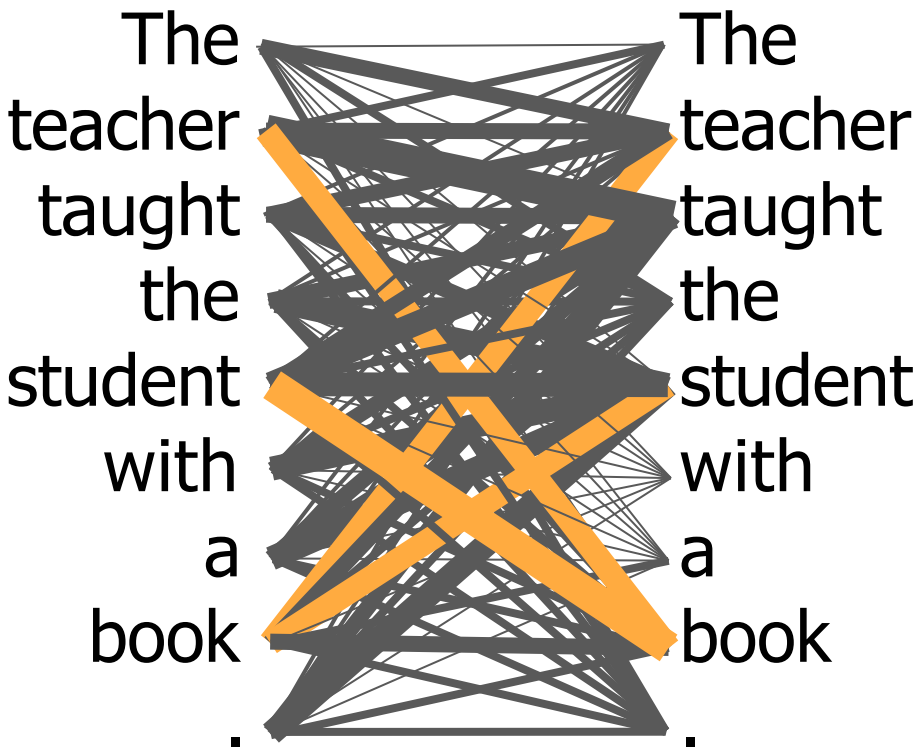

TRANSFORMERS

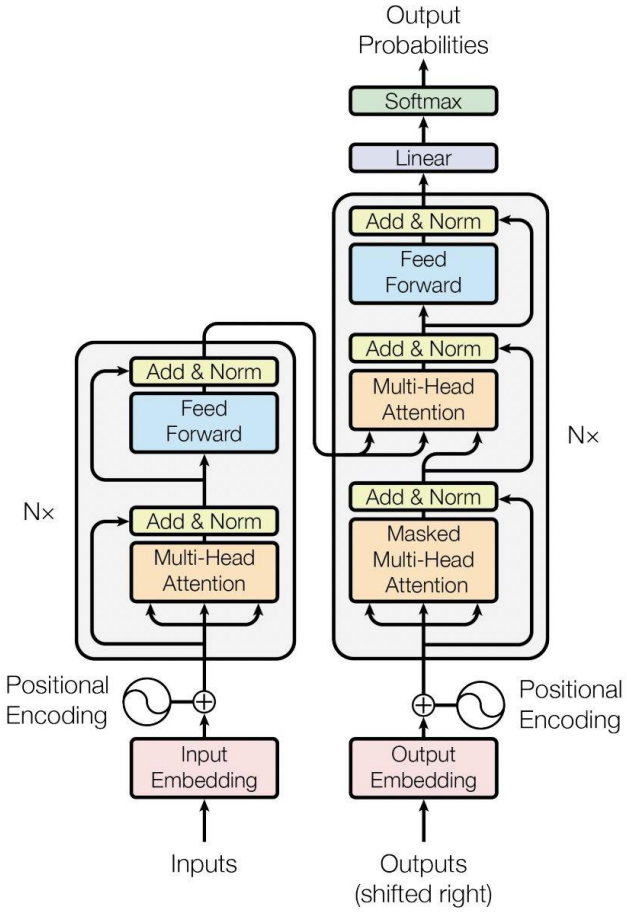

TRANSFORMERS

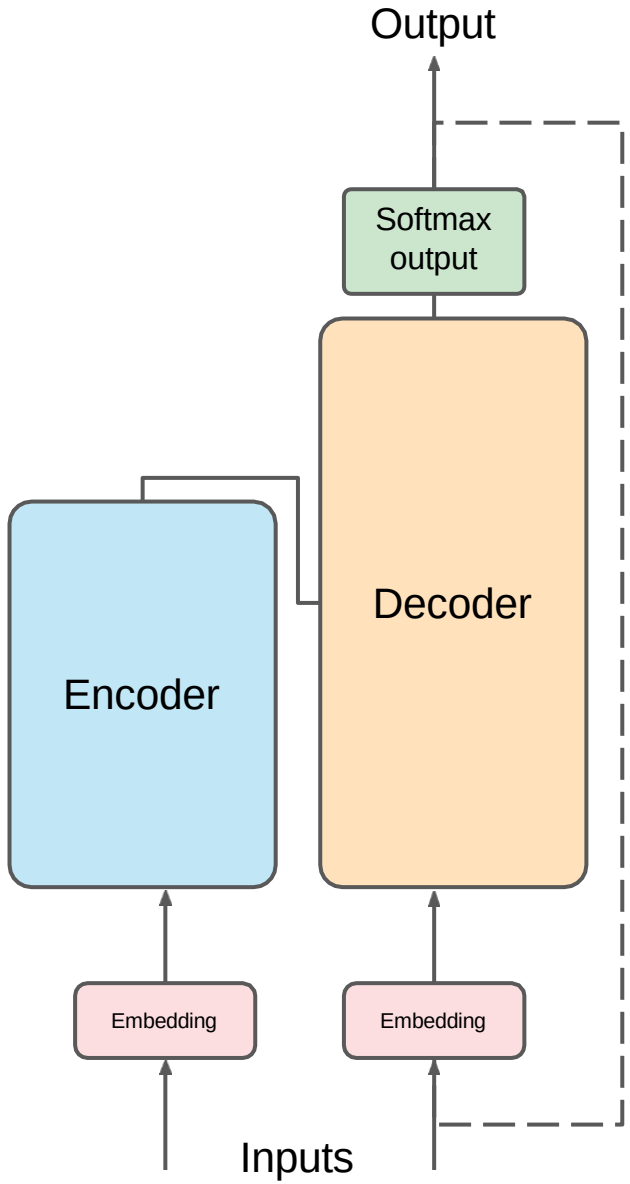

TRANSFORMERS

Encoder
Encodes inputs ("prompts") with contextual understanding and produces one vector per input token.

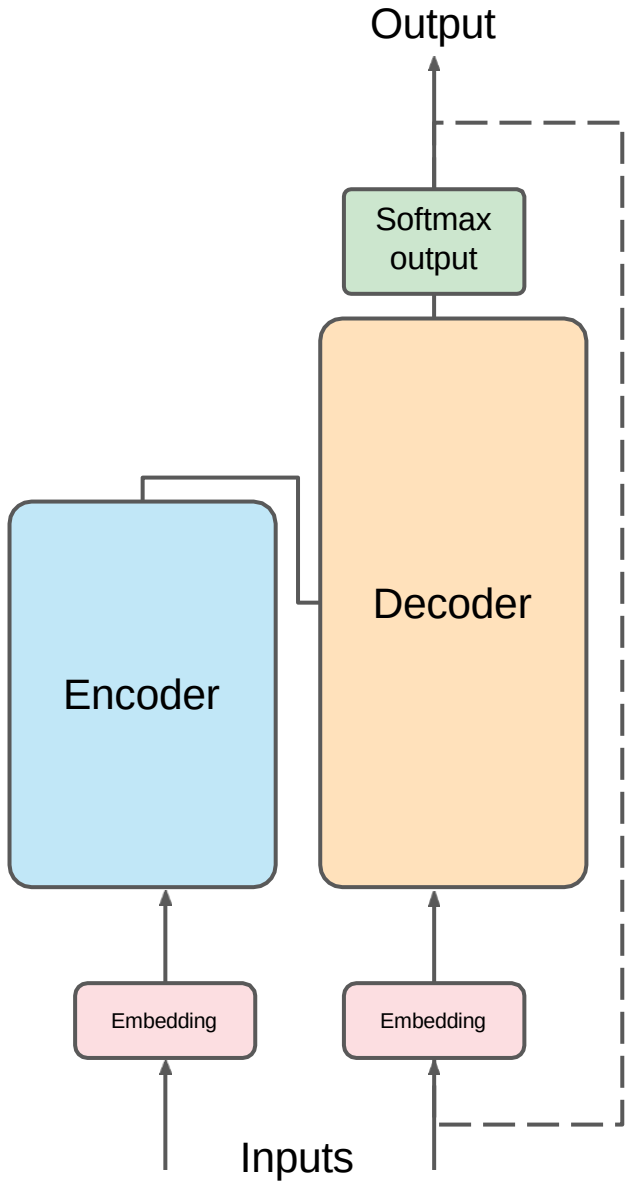

Decoder
Accepts input tokens and generates new tokens.

TRANSFORMERS

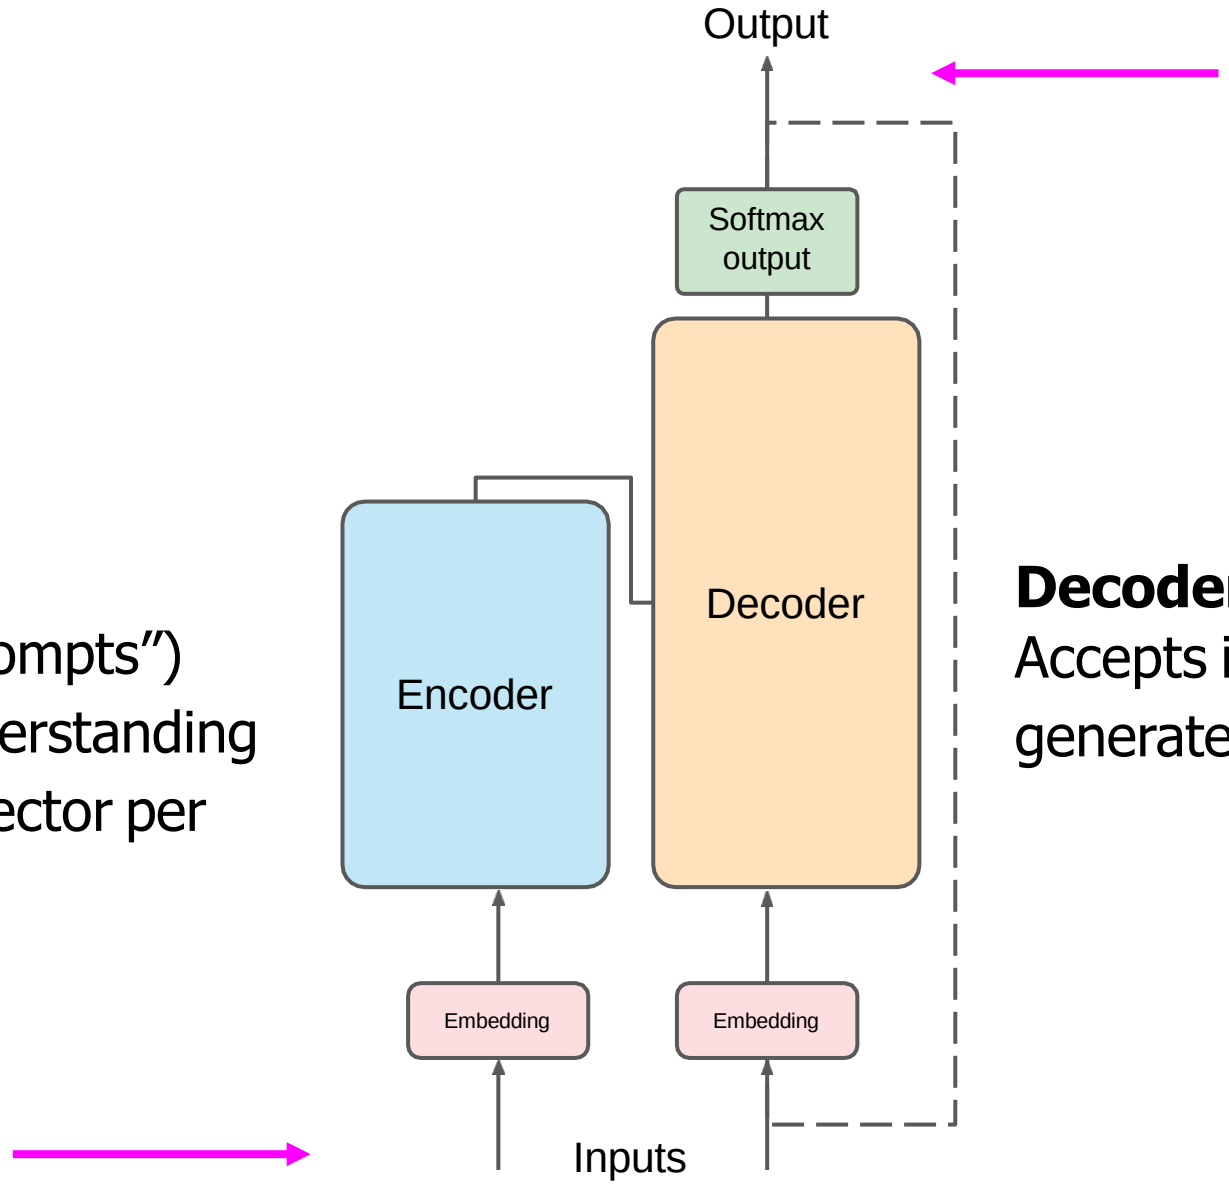

Encoder
Encodes inputs ("prompts") with contextual understanding and produces one vector per input token.

Decoder
Accepts input tokens and generates new tokens.

TRANSFORMERS

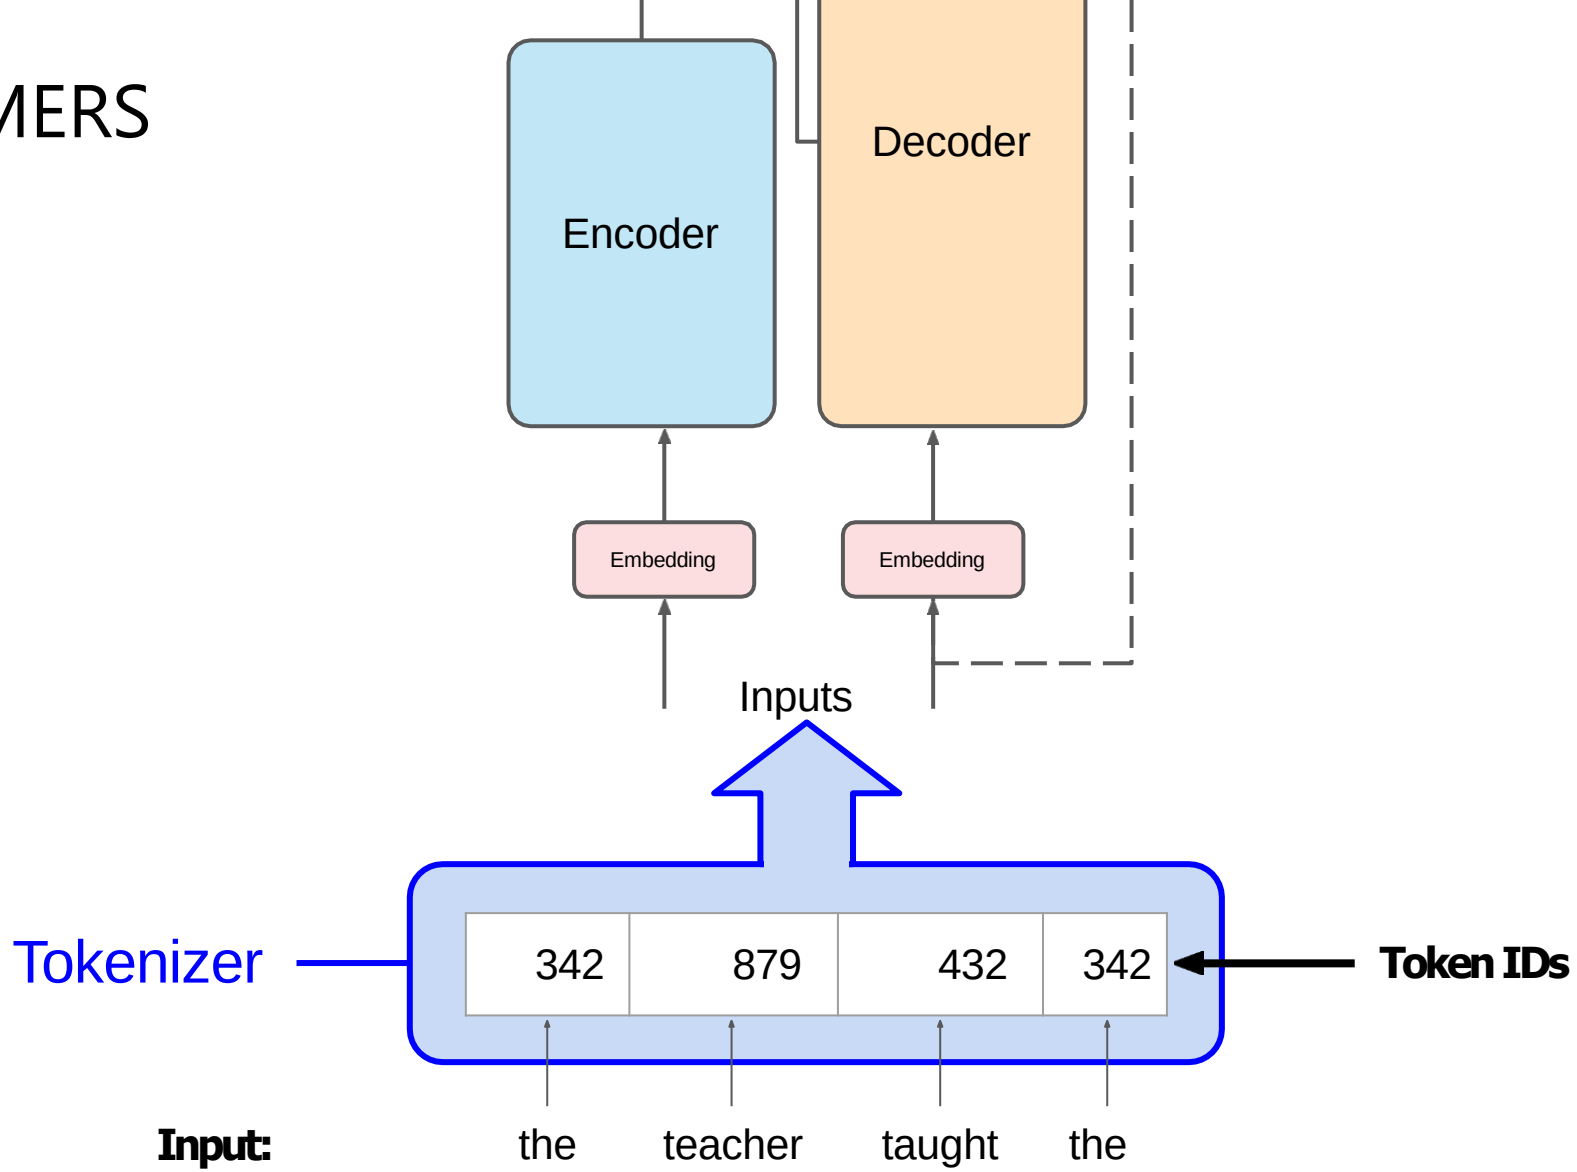

TRANSFORMERS

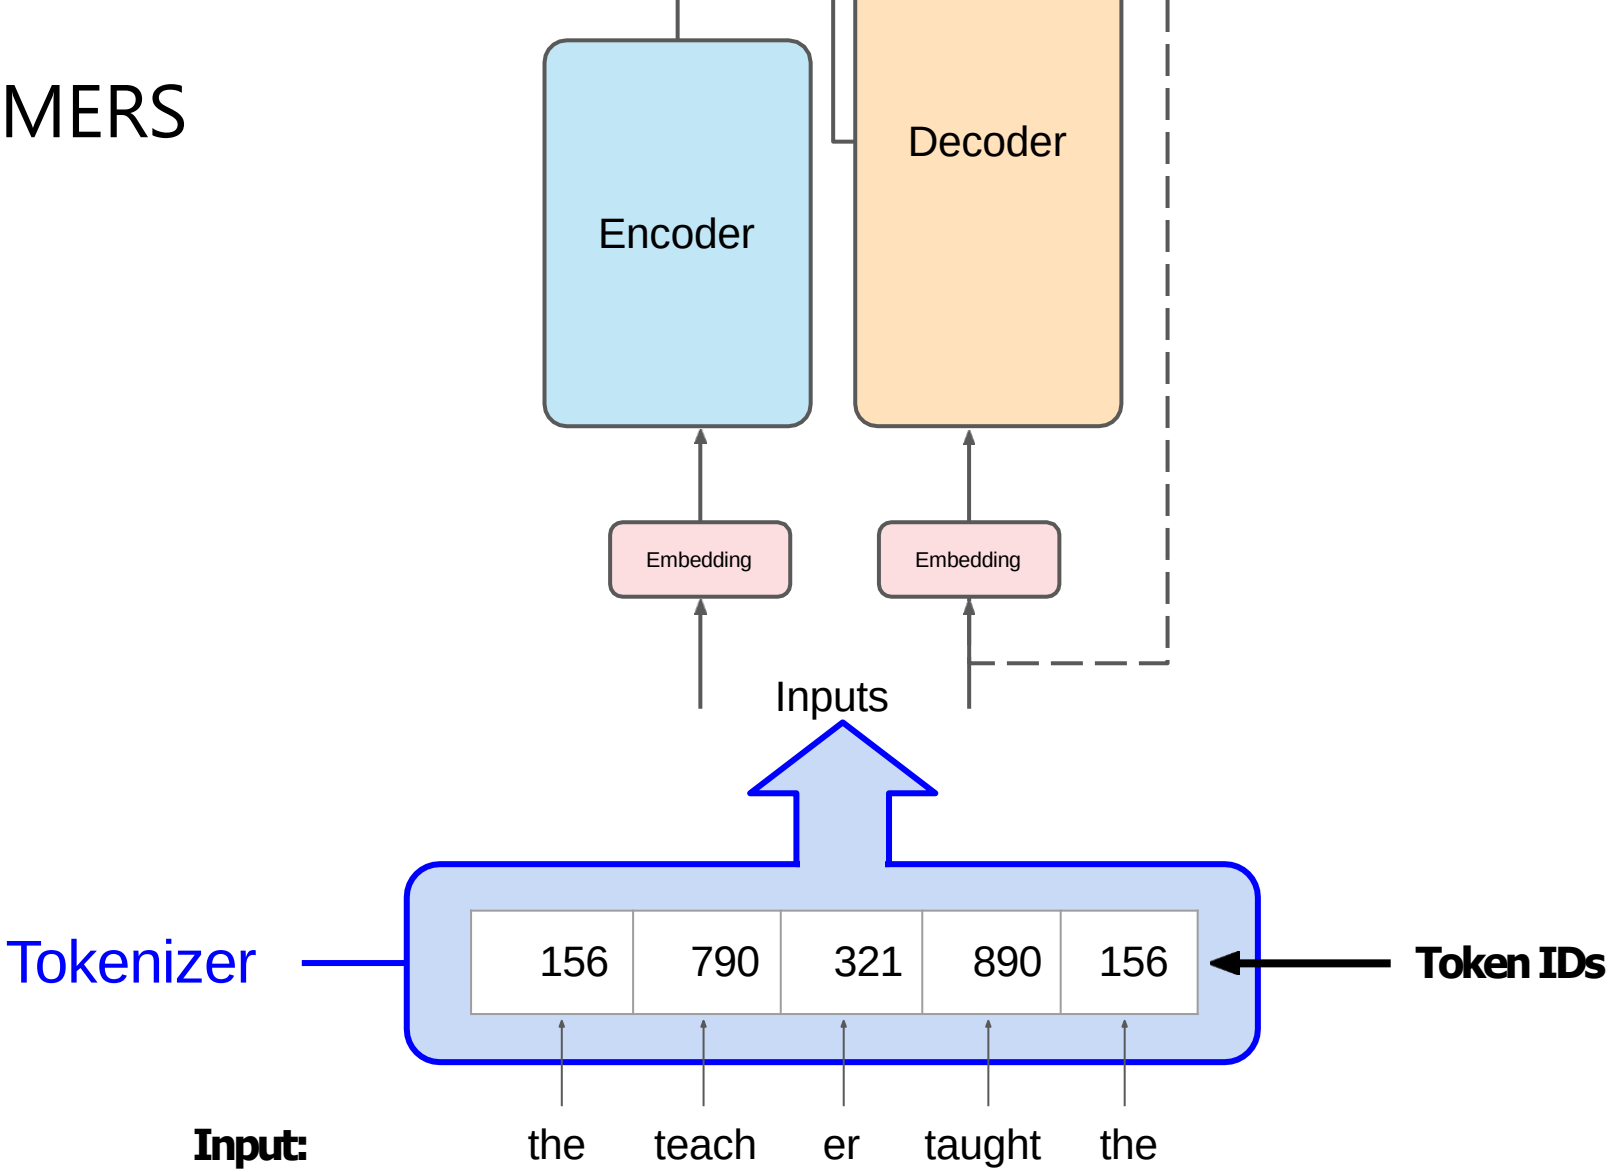

TRANSFORMERS

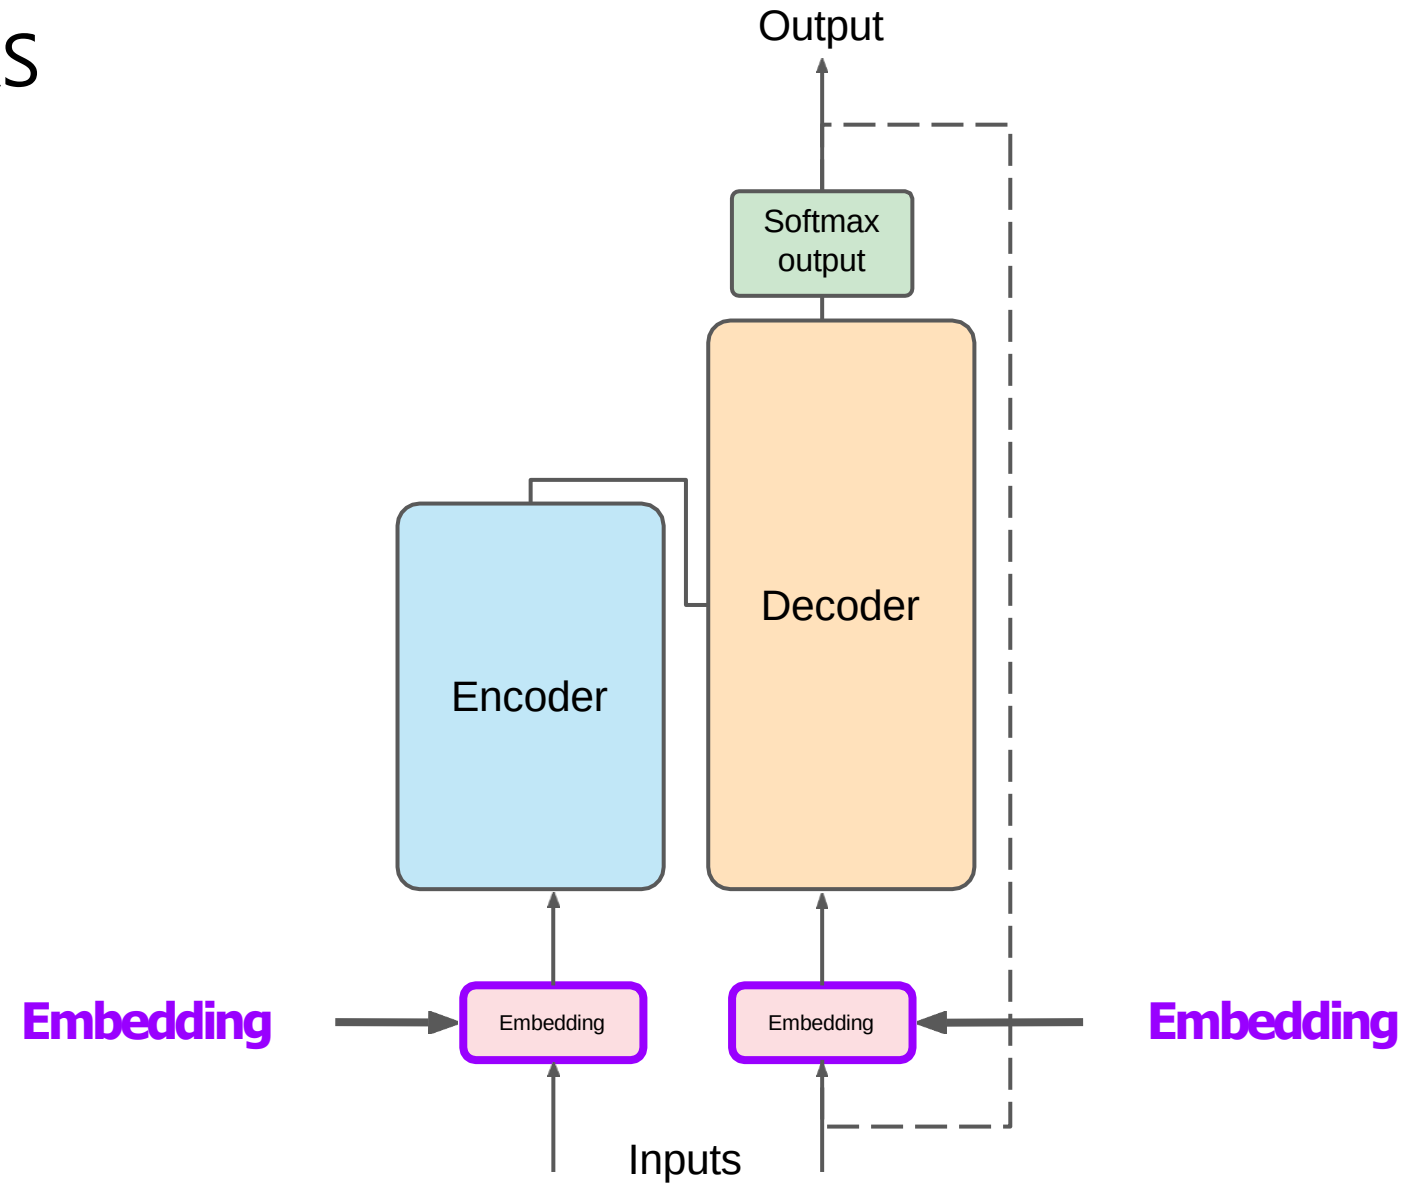

TRANSFORMERS

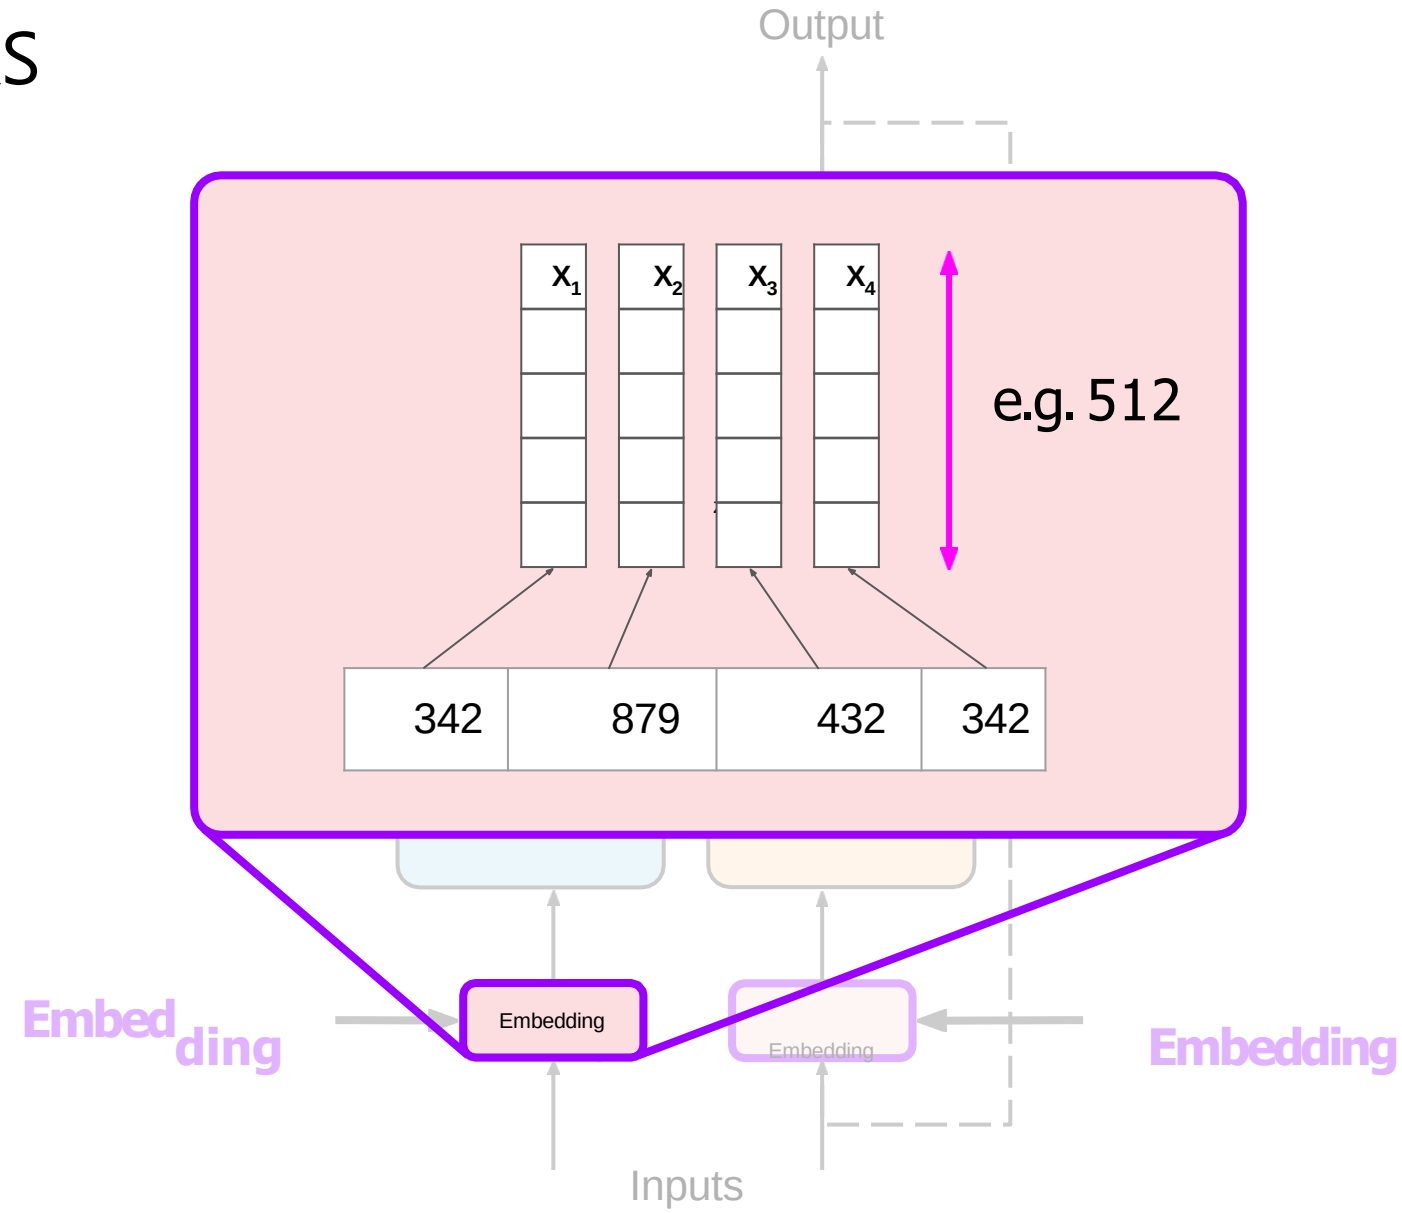

TRANSFORMERS

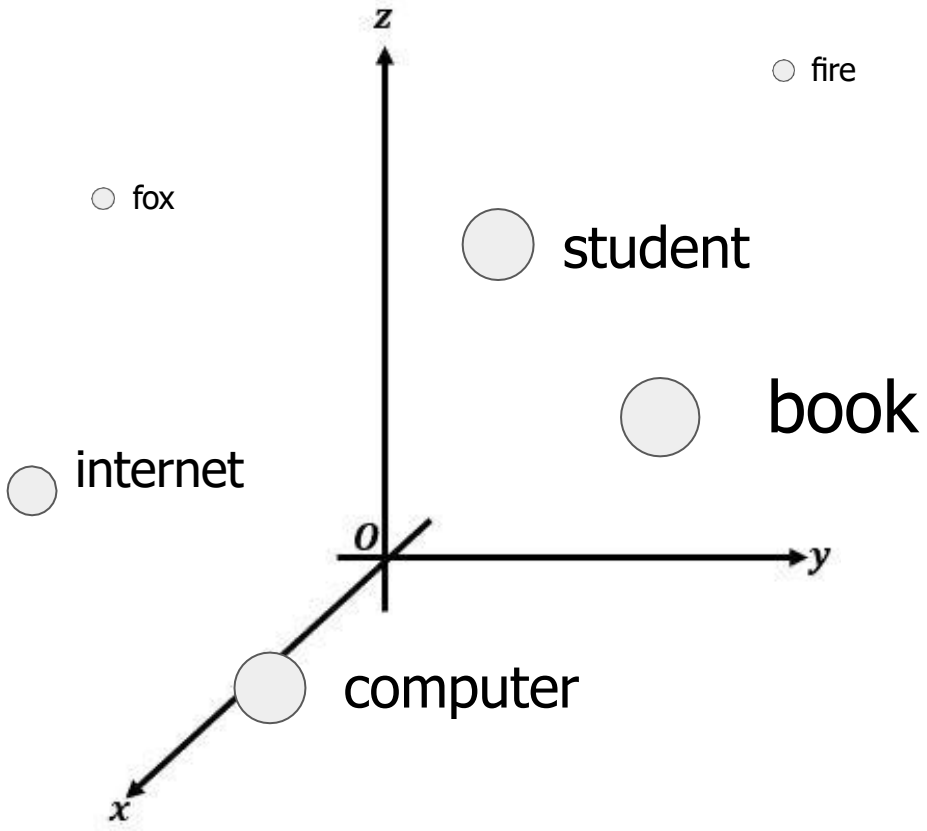

TRANSFORMERS

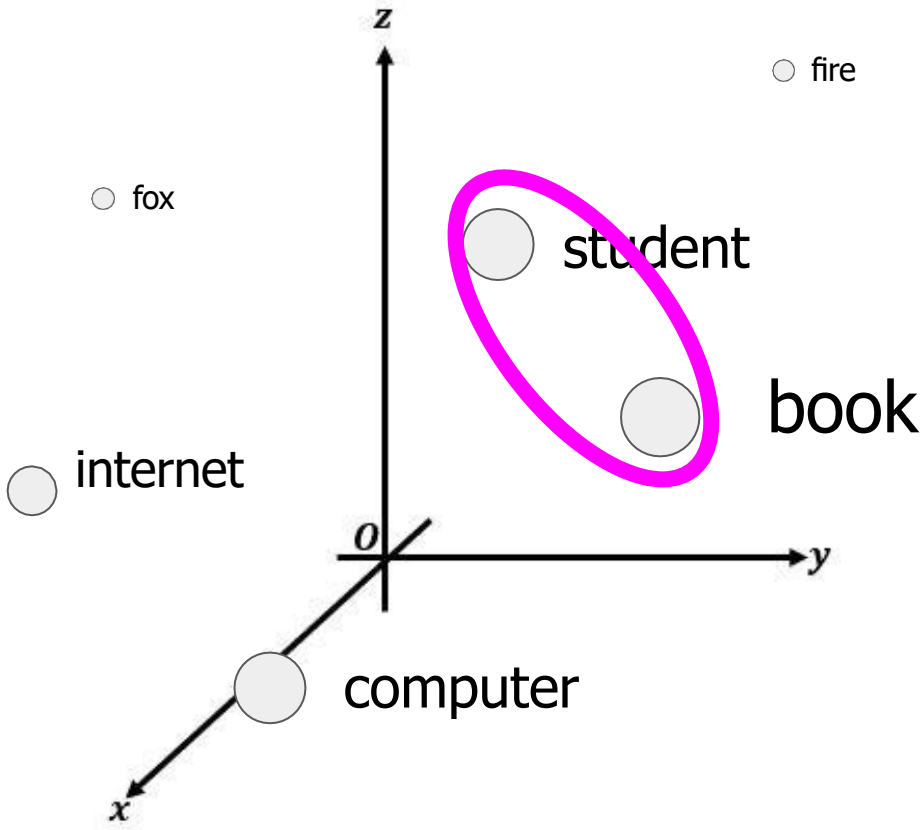

TRANSFORMERS

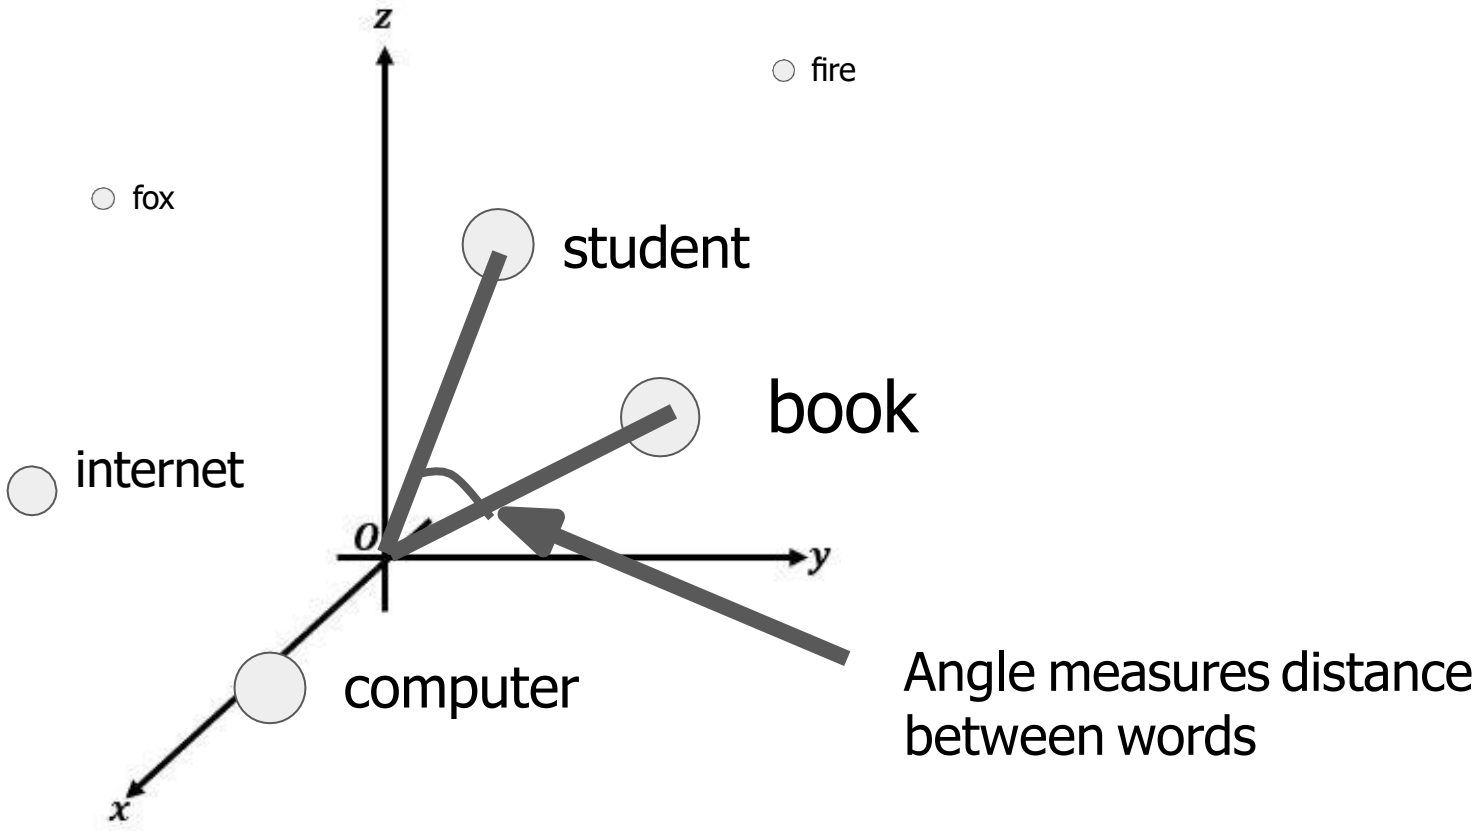

TRANSFORMERS

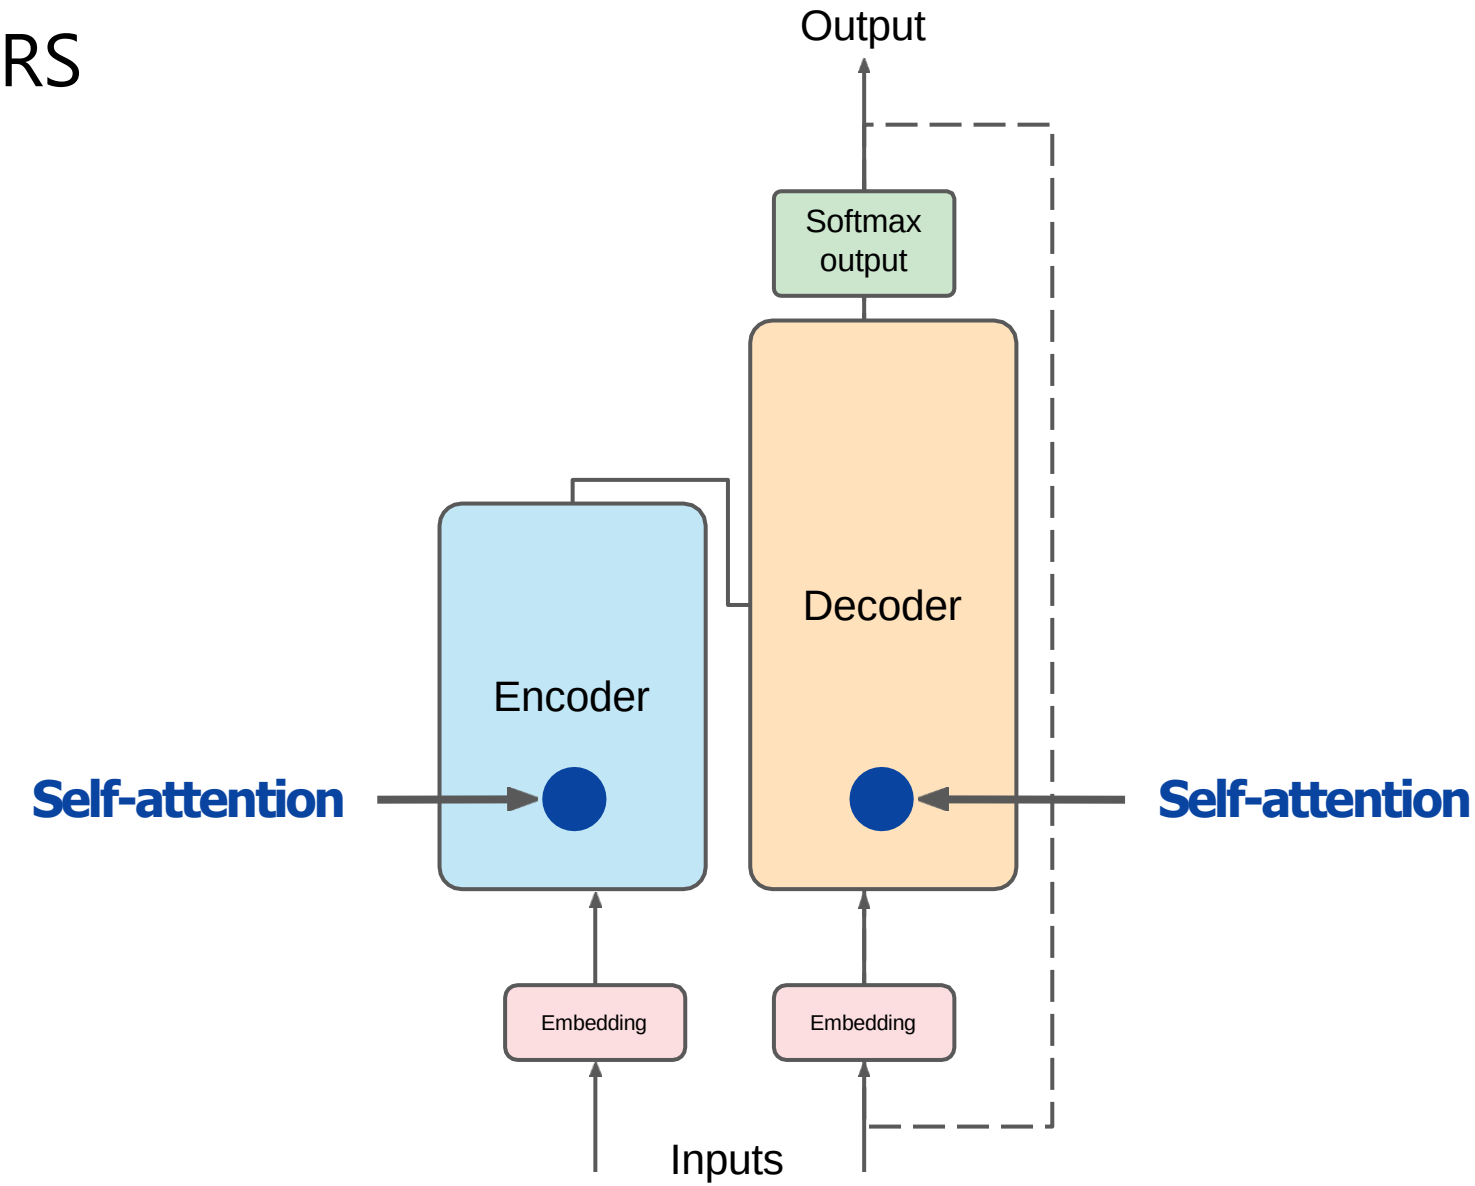

TRANSFORMERS

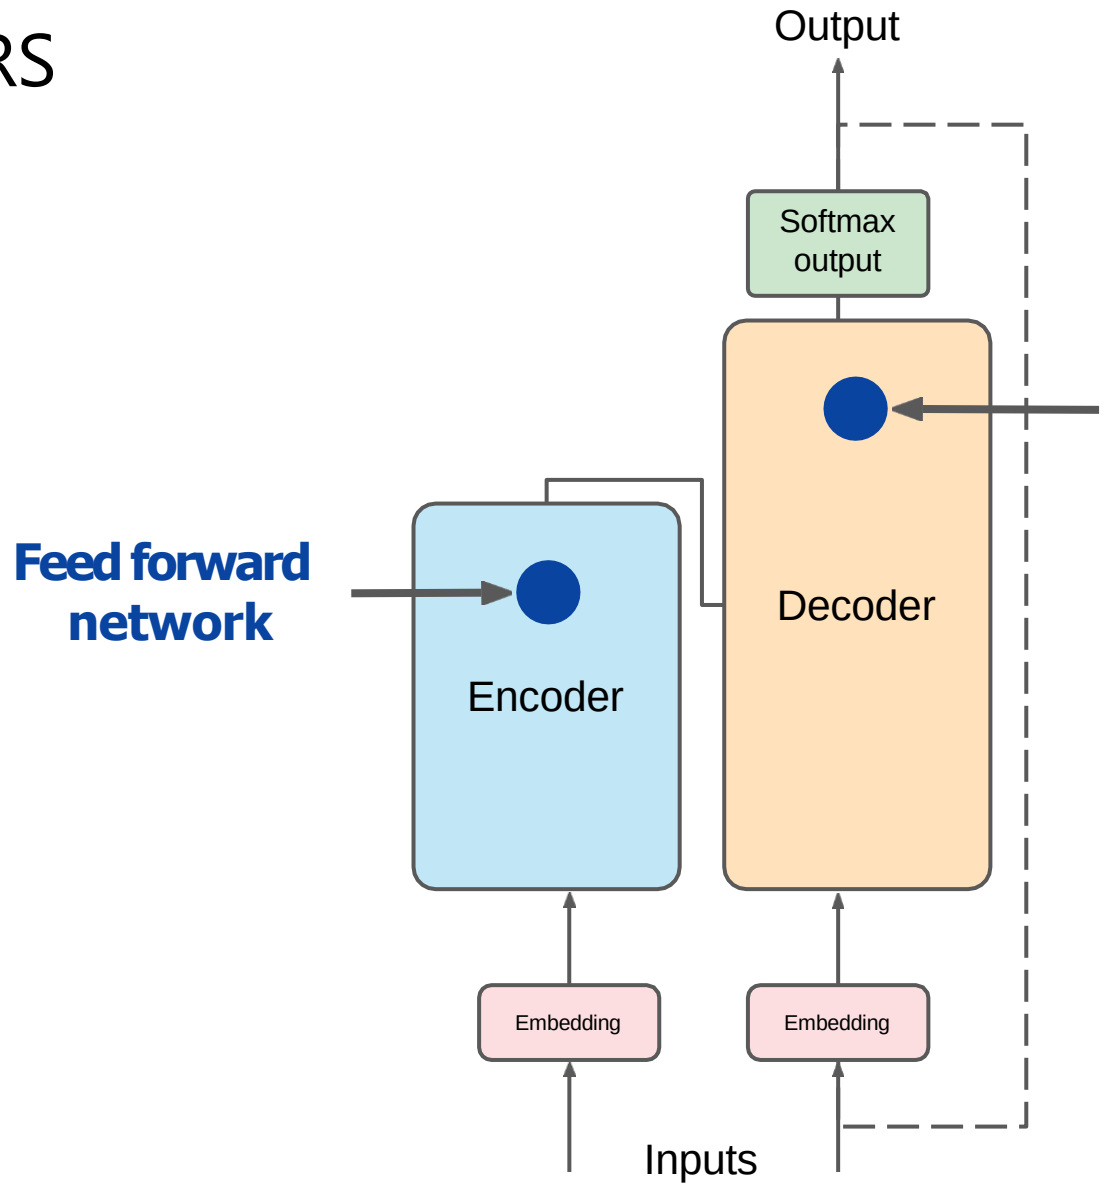

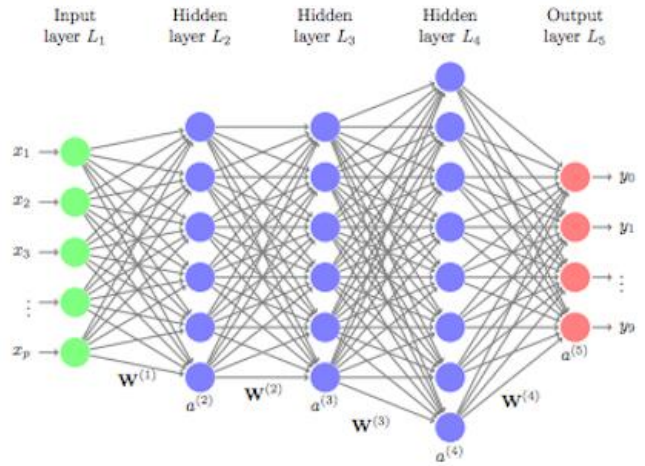

TRANSFORMERS

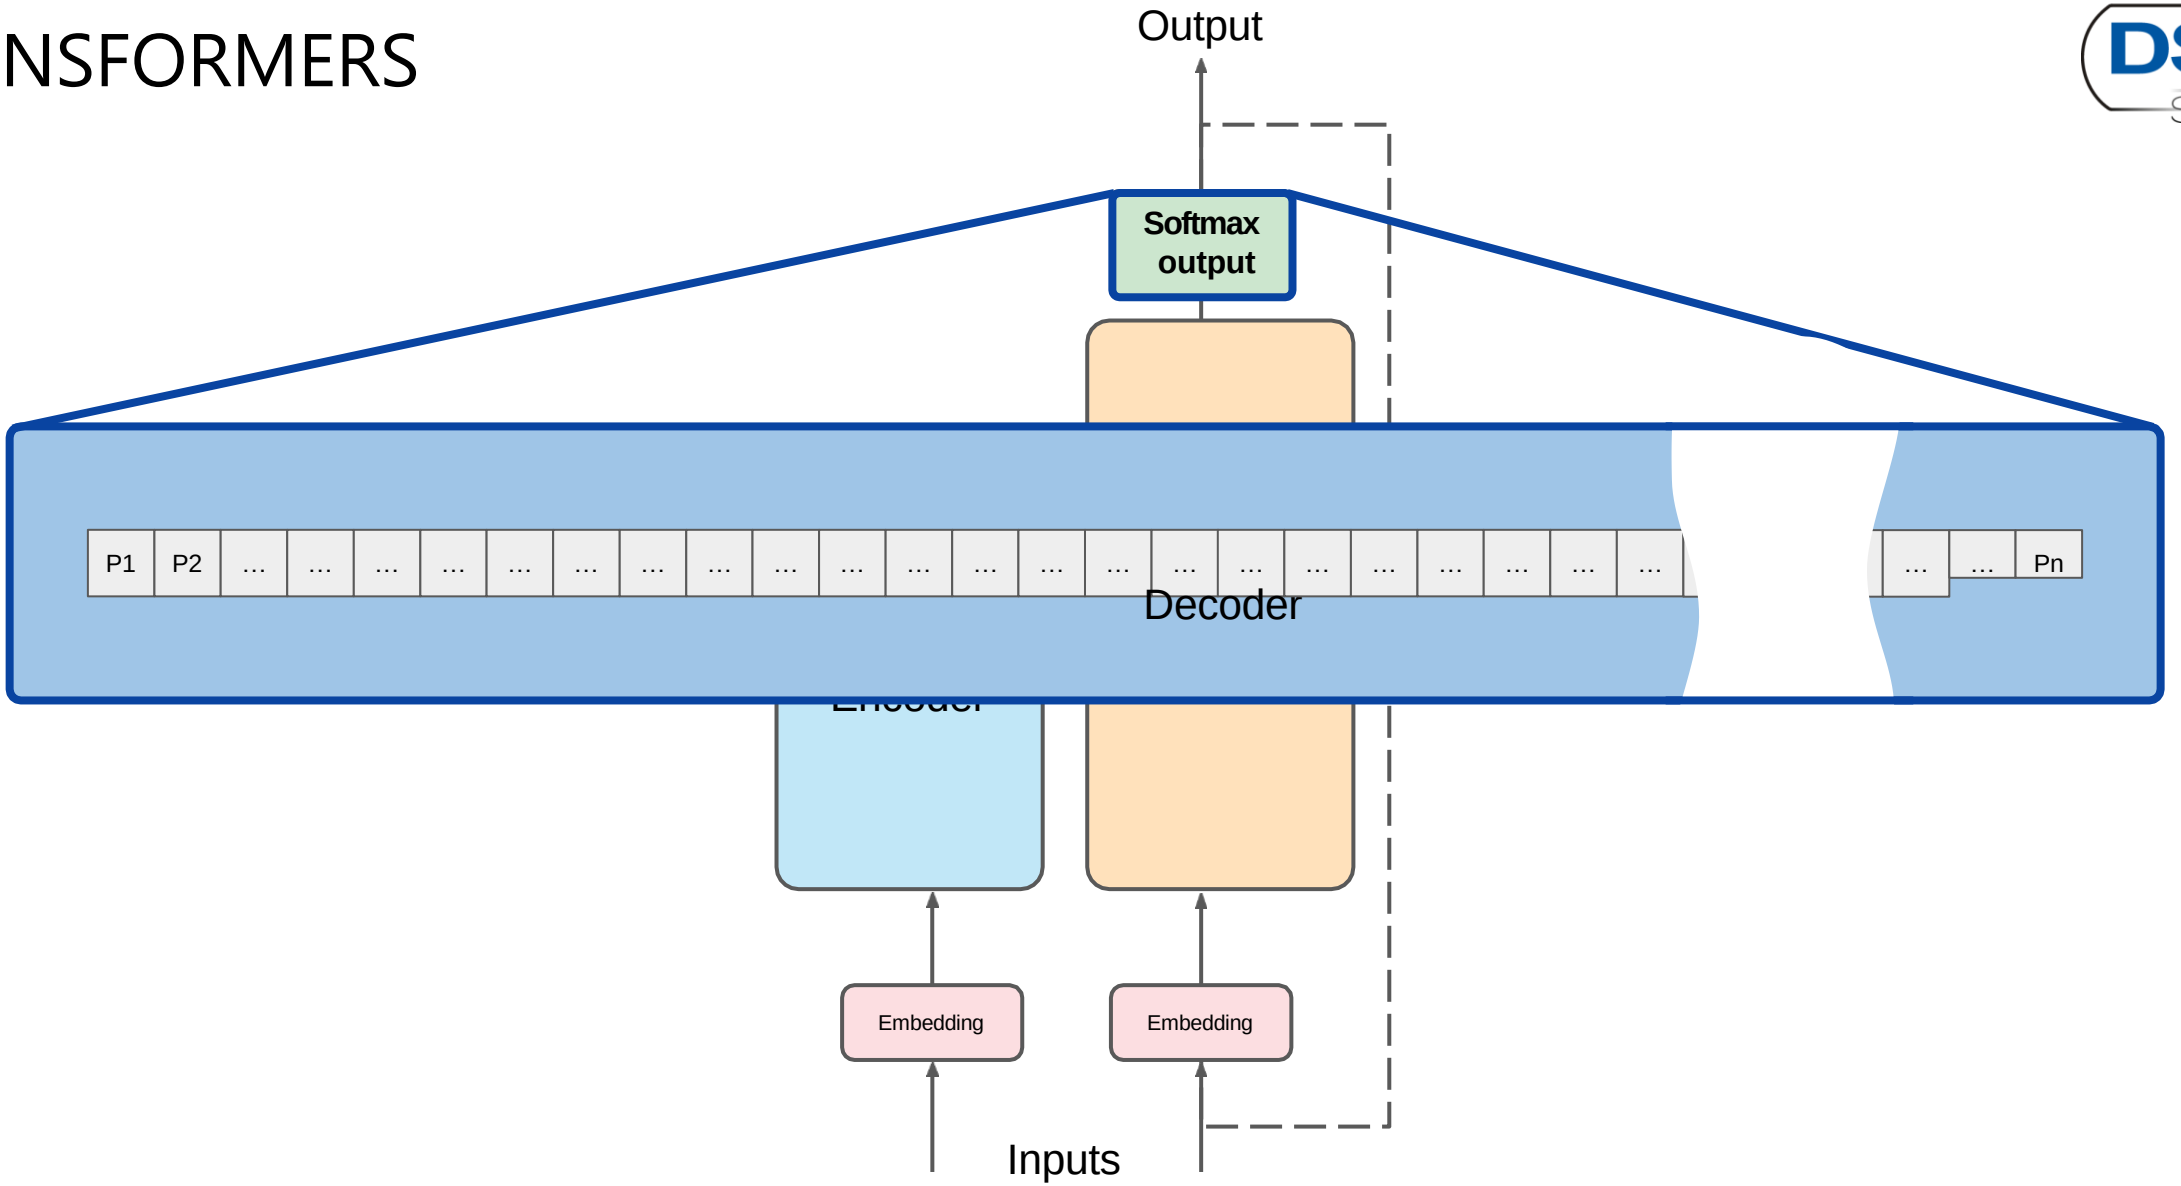

TRANSFORMERS

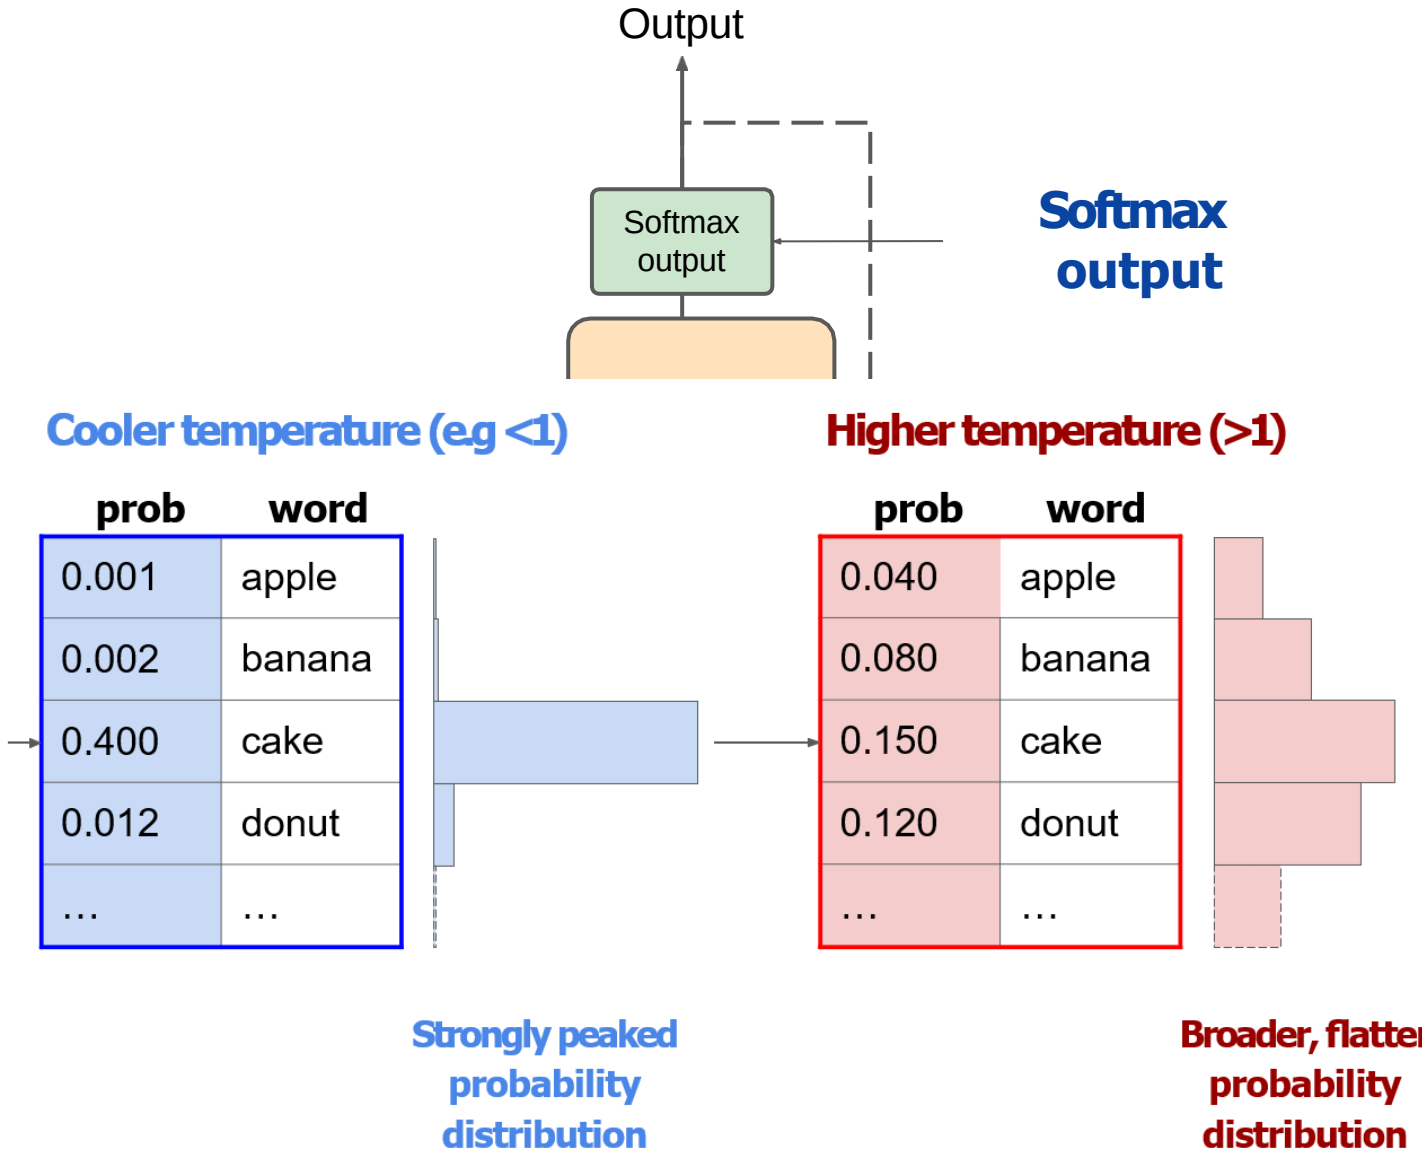

TRANSFORMERS

Translation: sequence-to-sequence task

`Milujem seminár DSA`

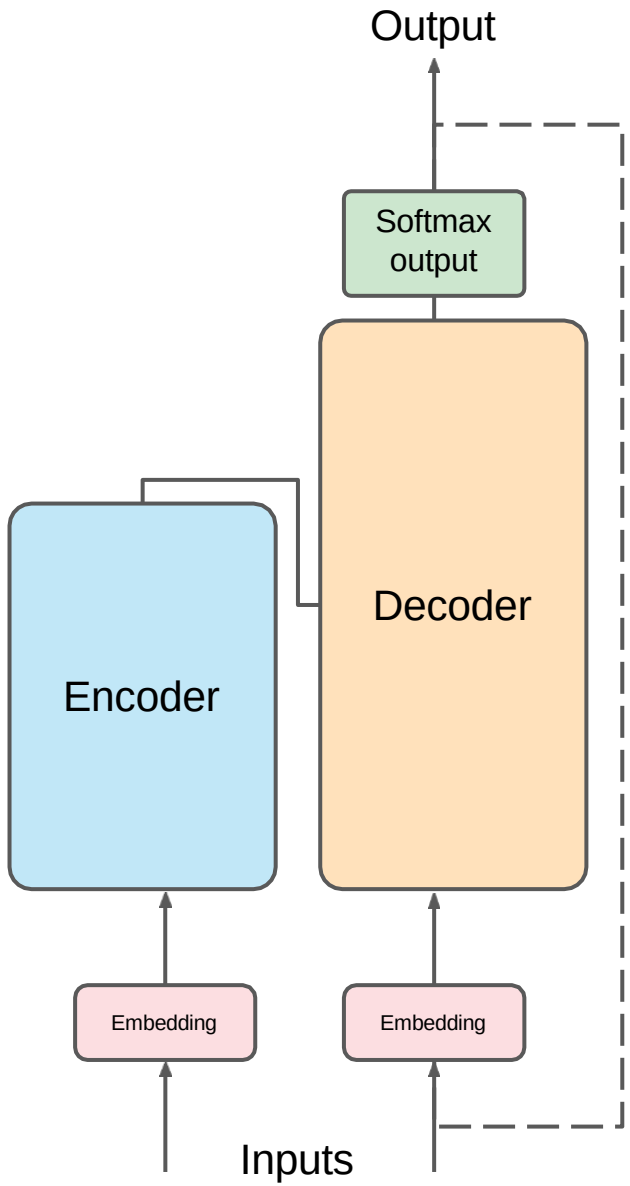

TRANSFORMERS

Translation: sequence-to-sequence task

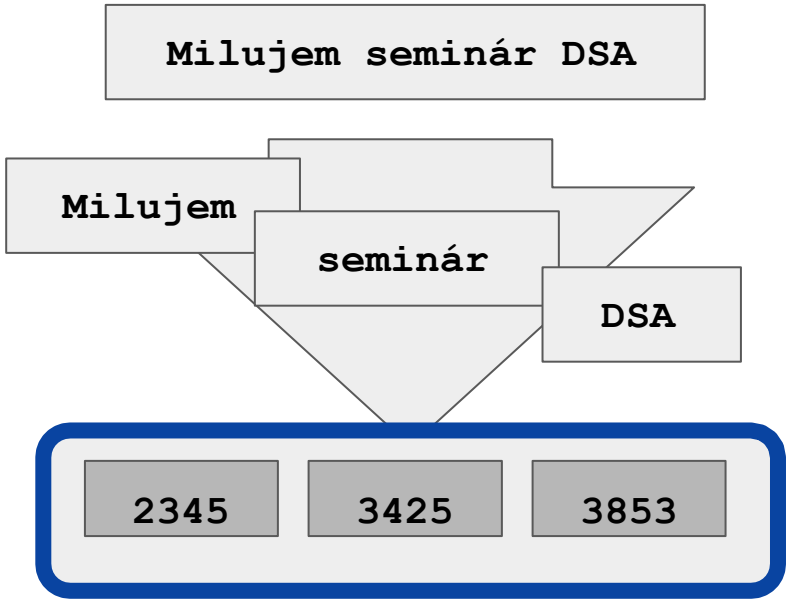

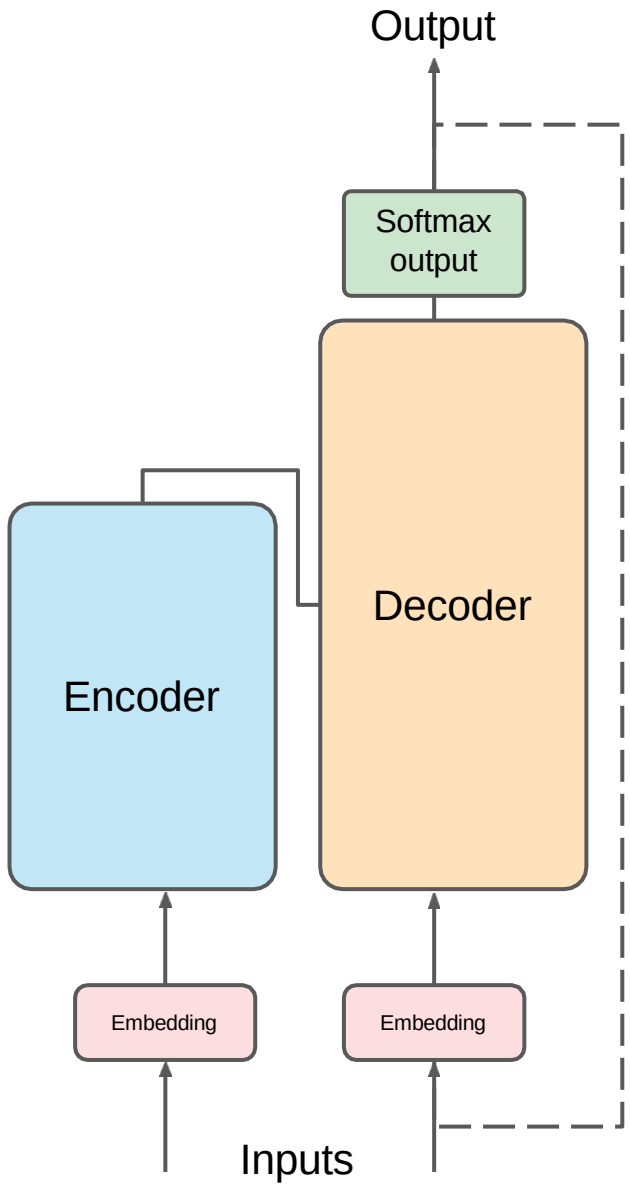

TRANSFORMERS

Translation: sequence-to-sequence task

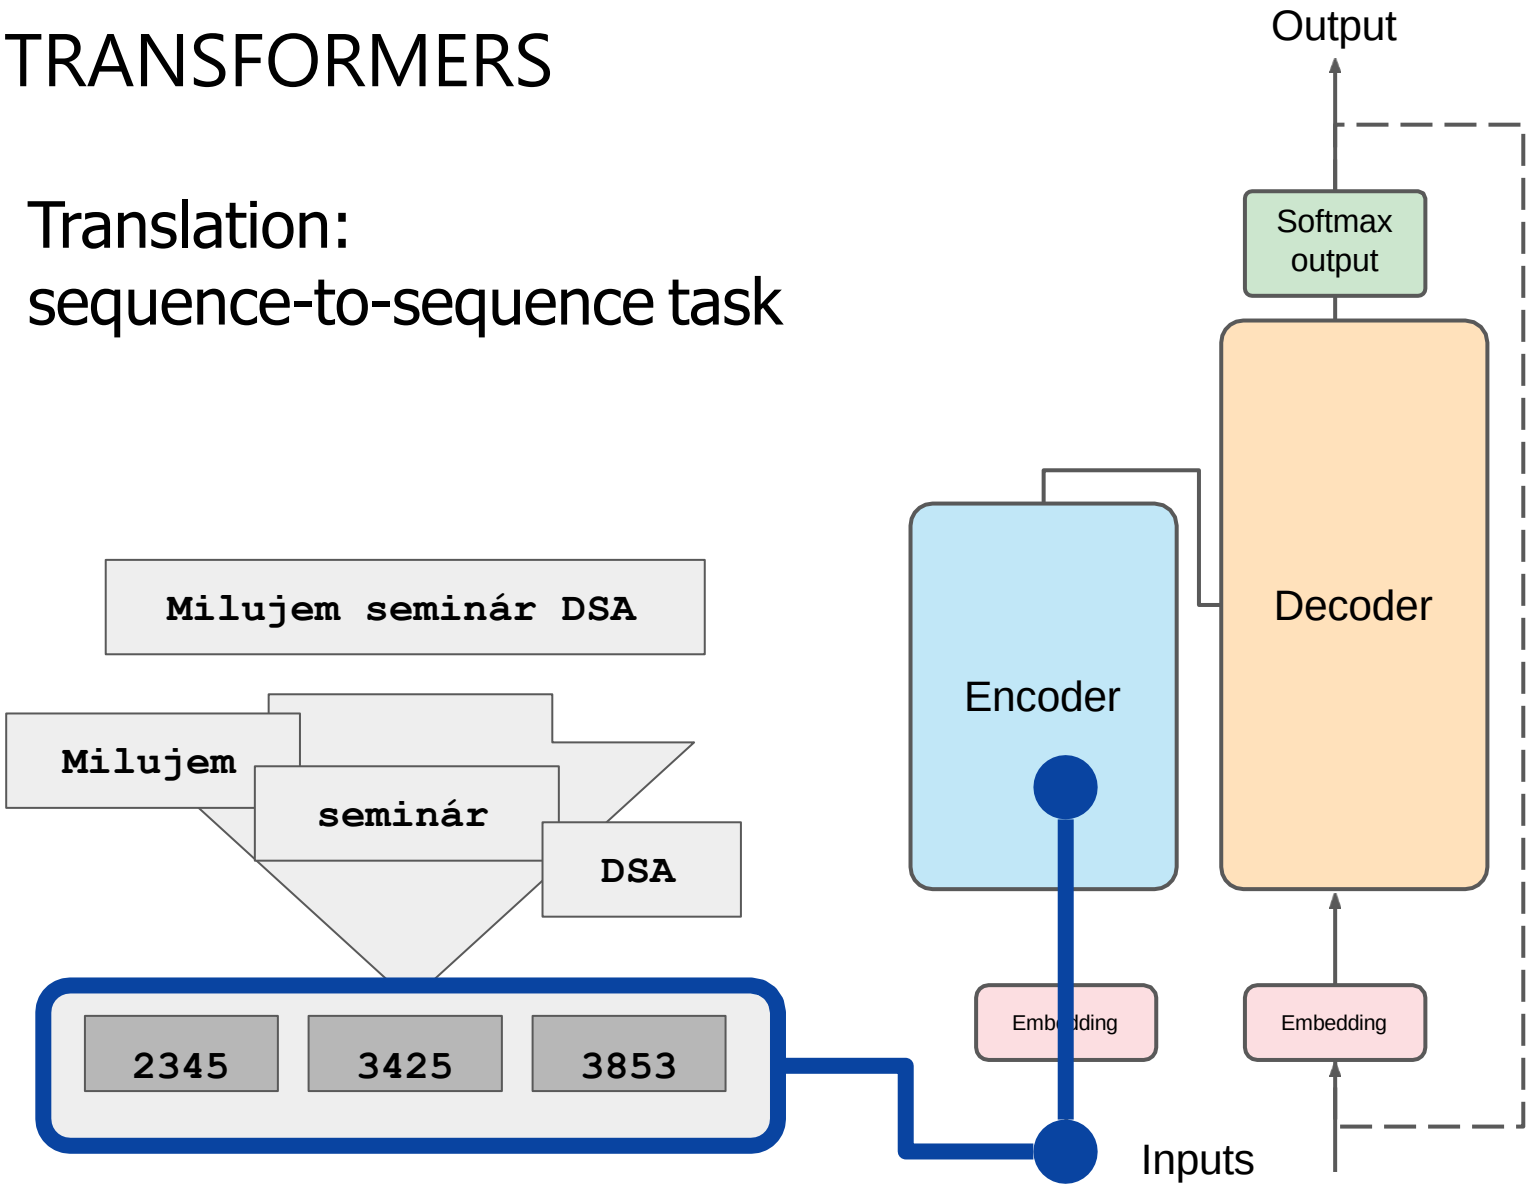

TRANSFORMERS

Translation: sequence-to-sequence task

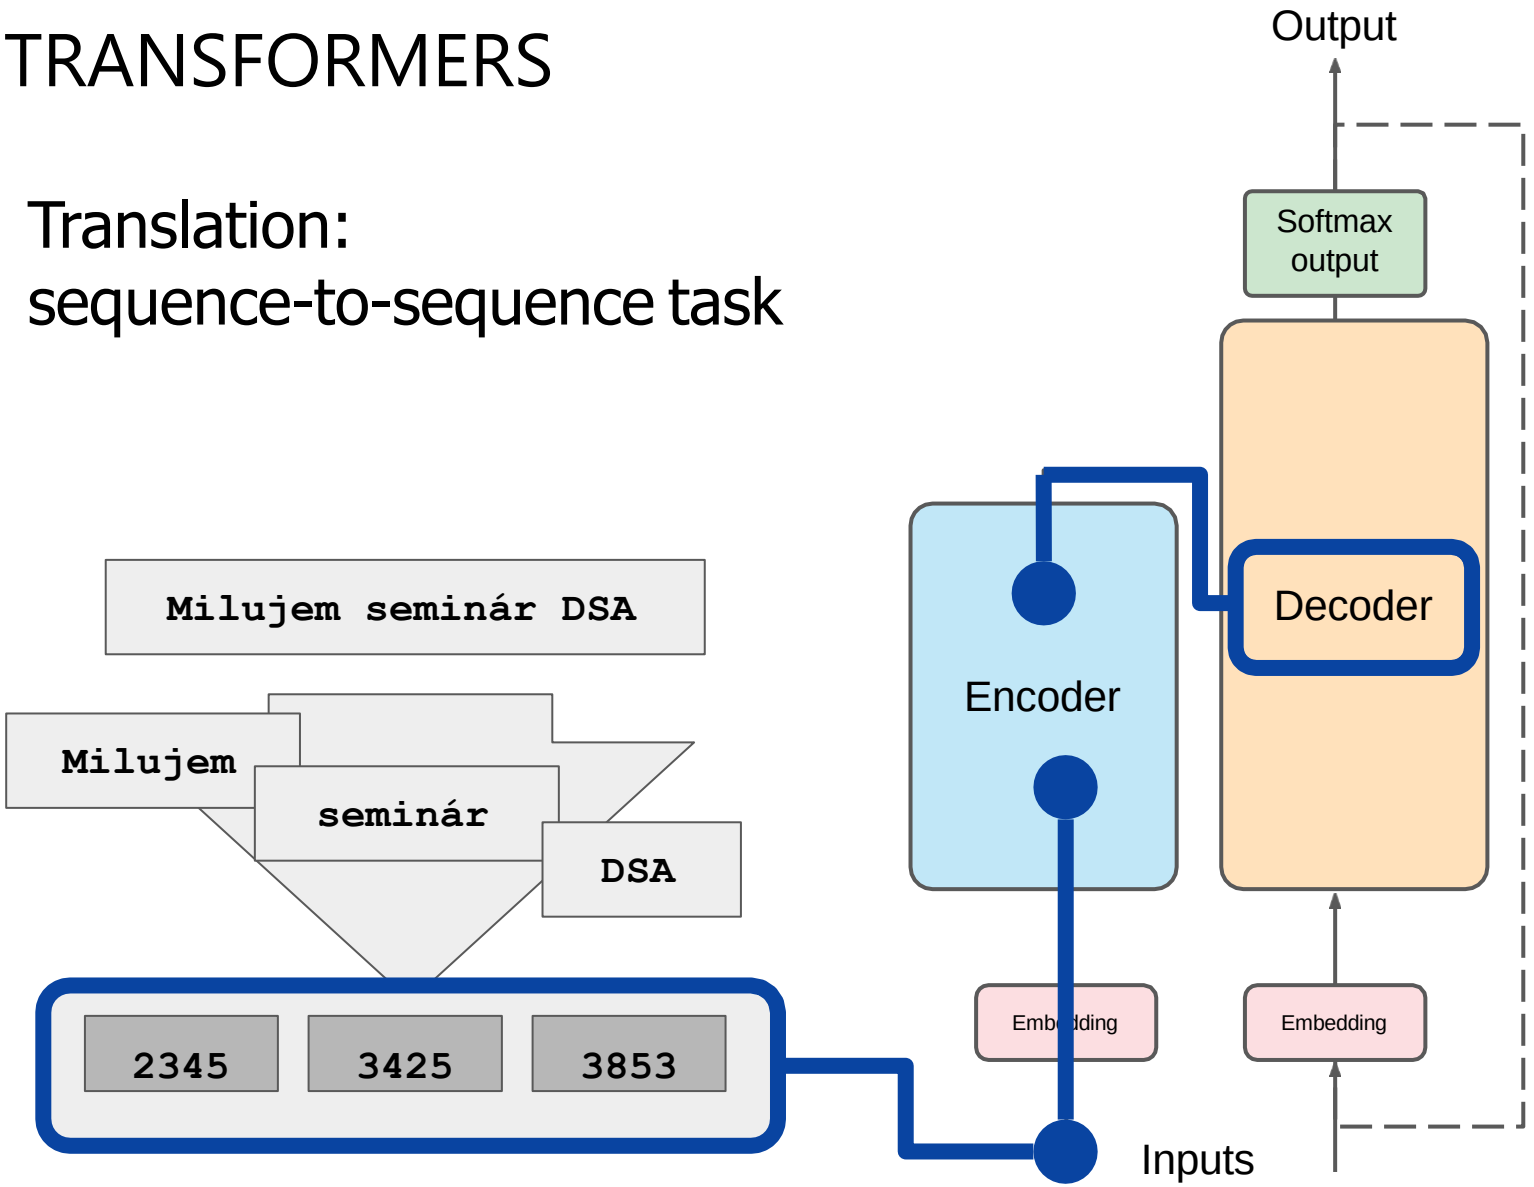

TRANSFORMERS

Translation: sequence-to-sequence task

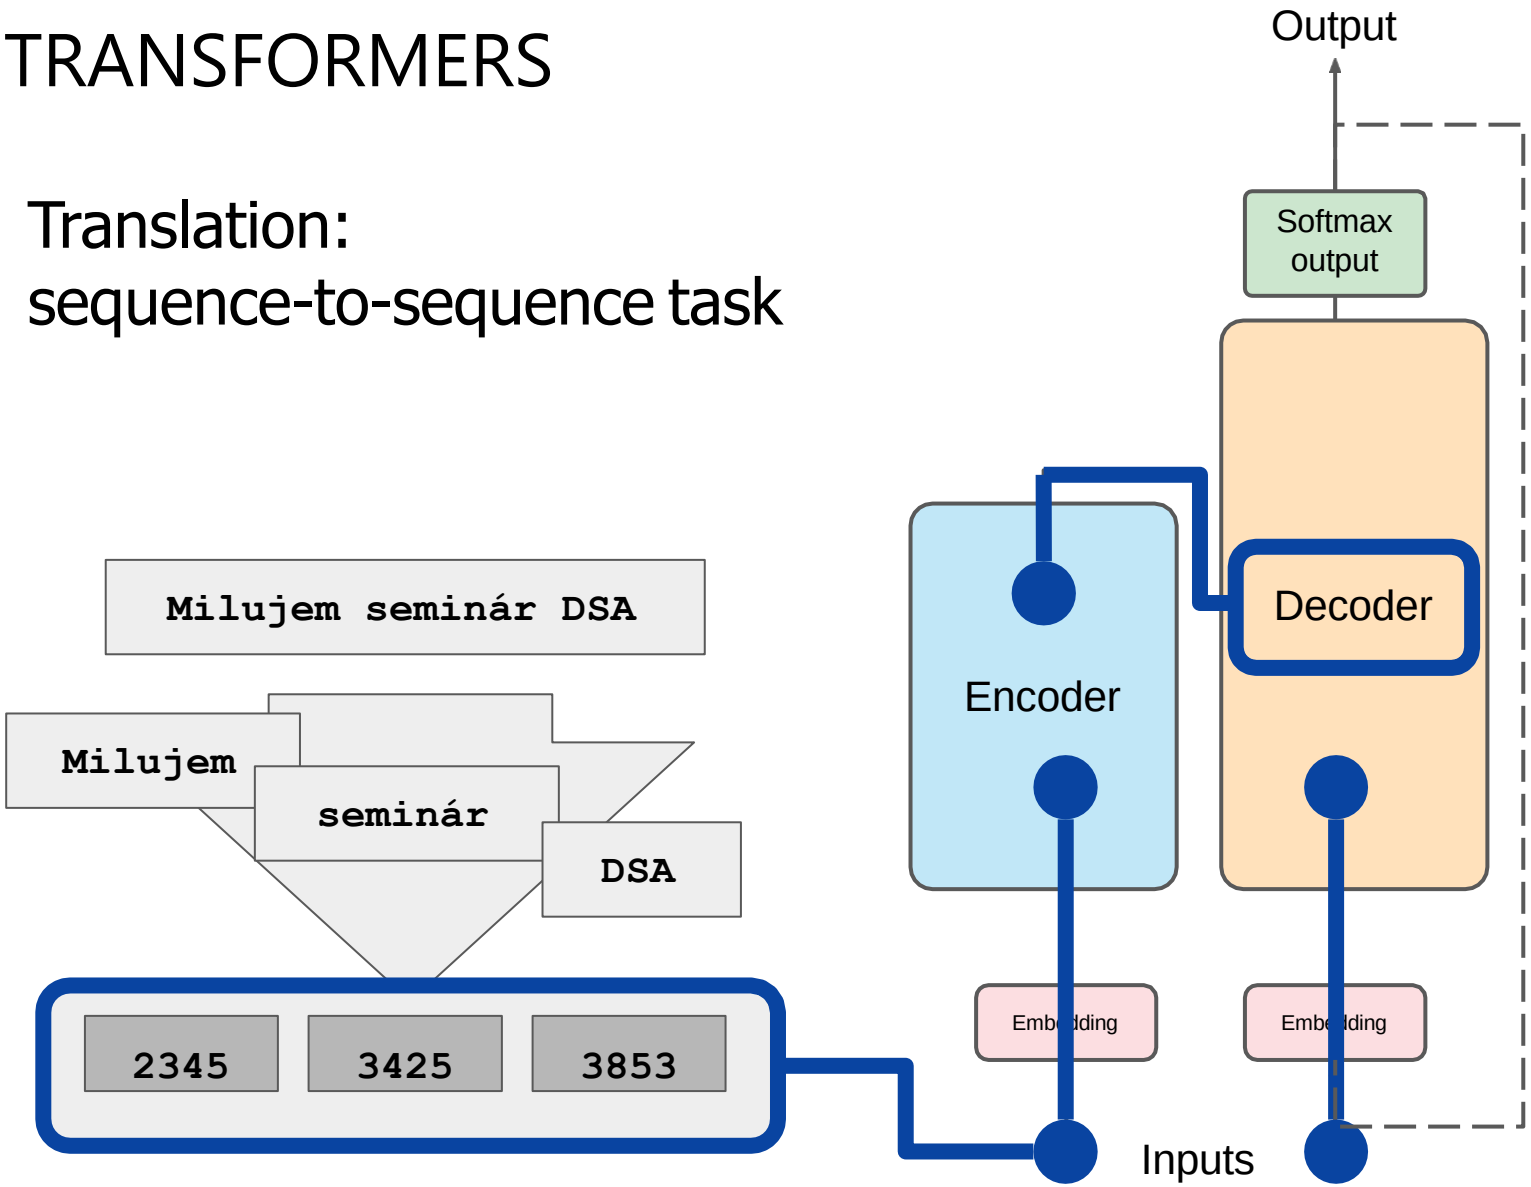

TRANSFORMERS

Translation: sequence-to-sequence task

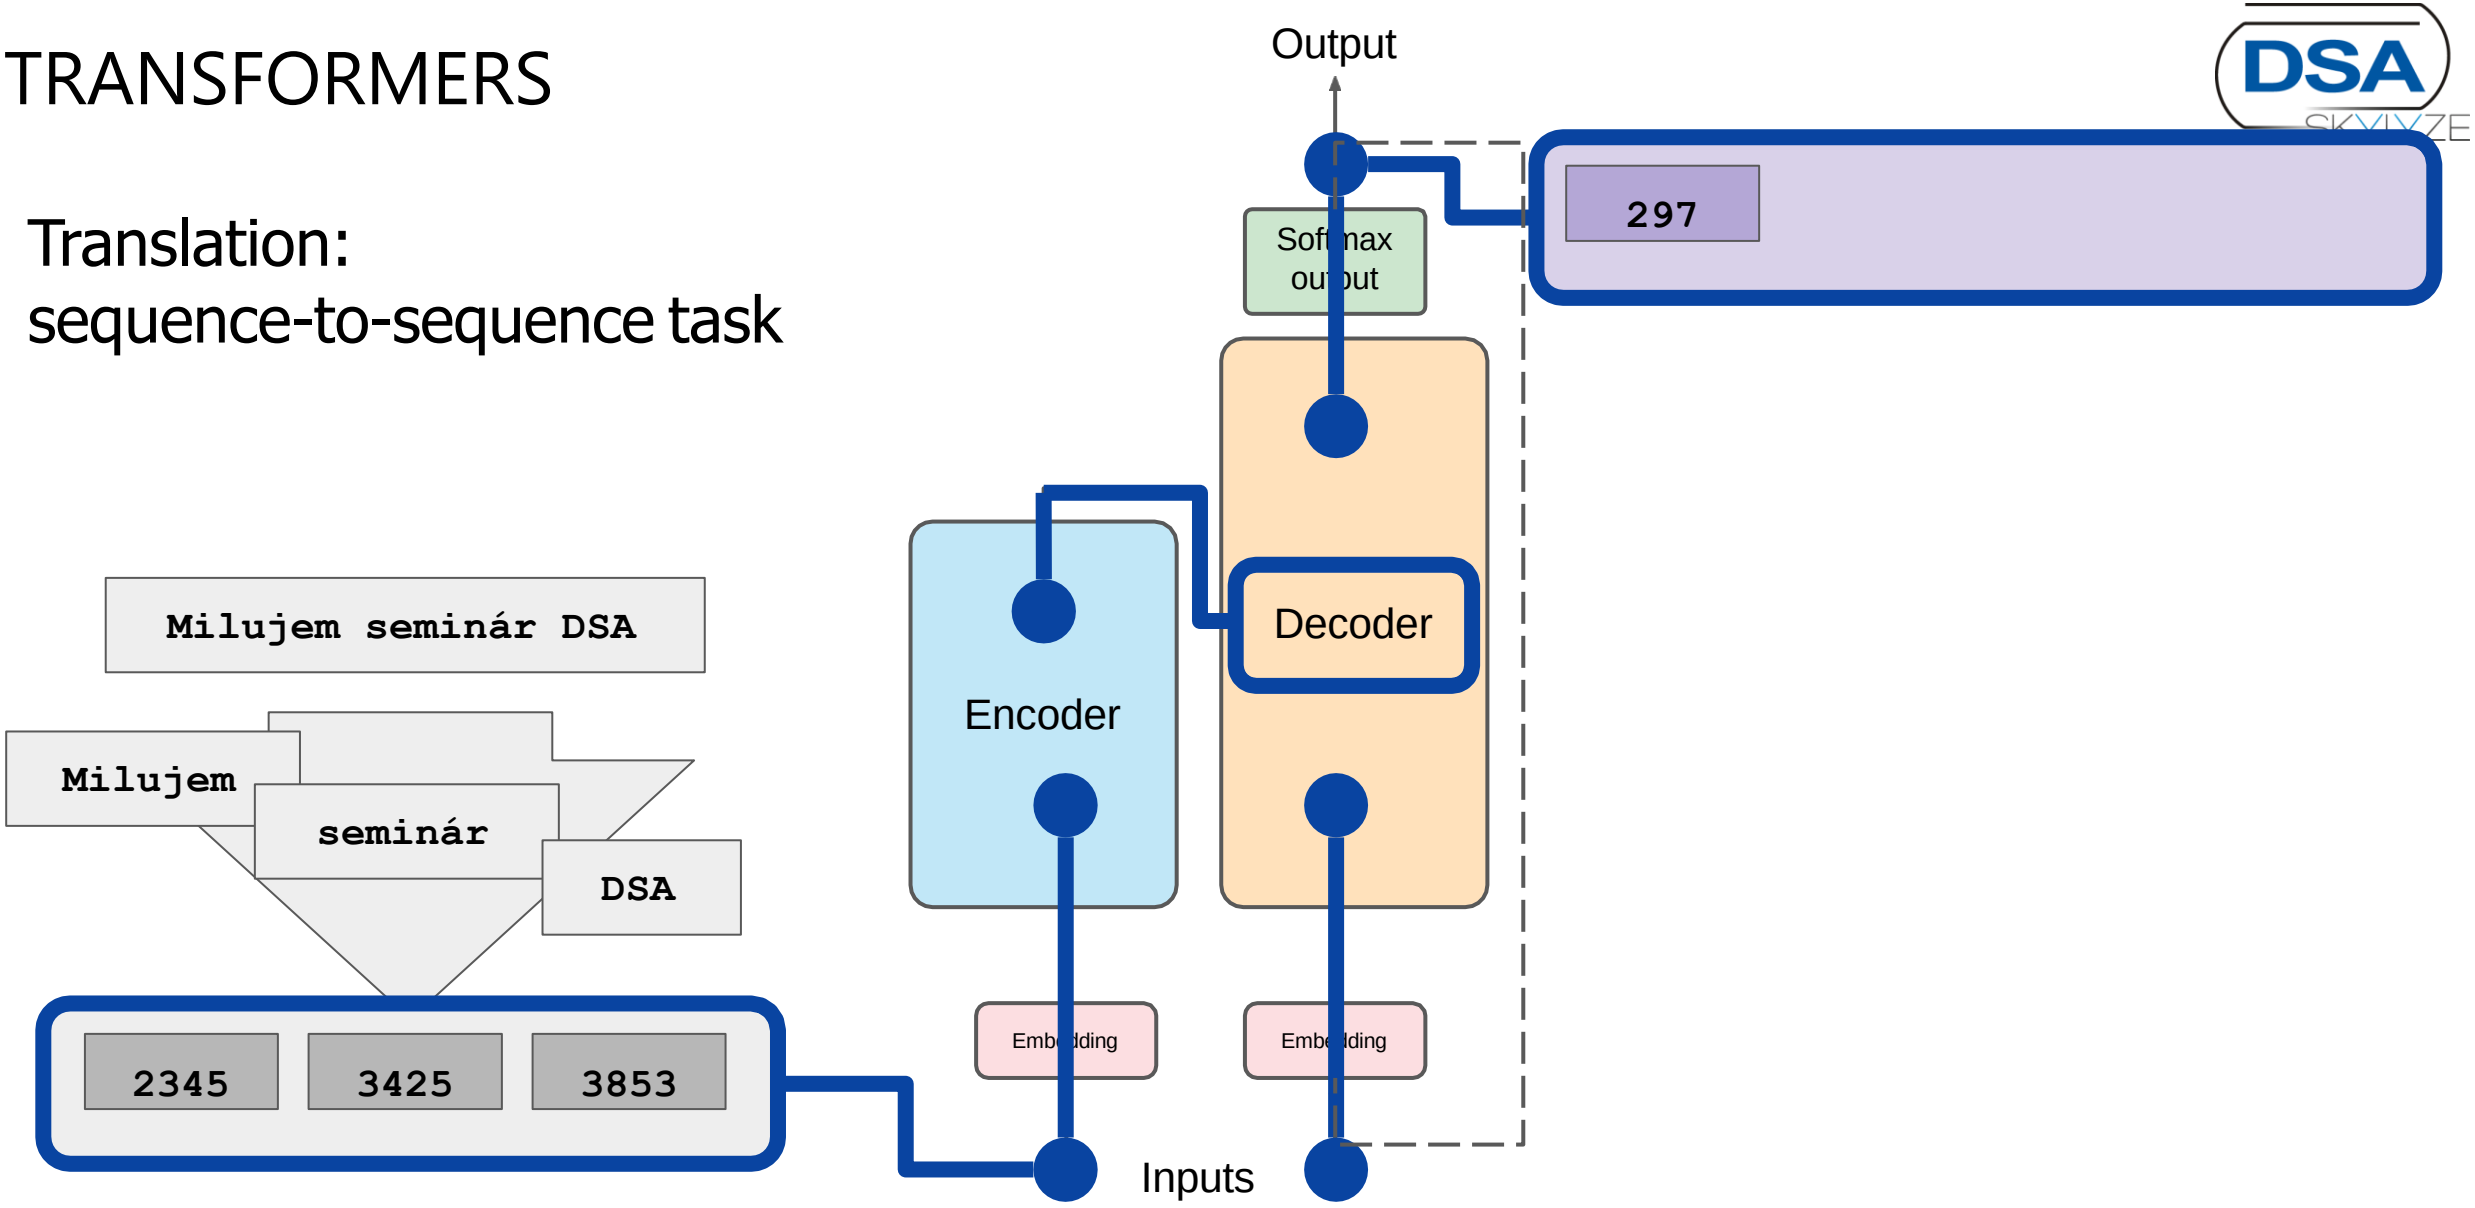

TRANSFORMERS

Translation: sequence-to-sequence task

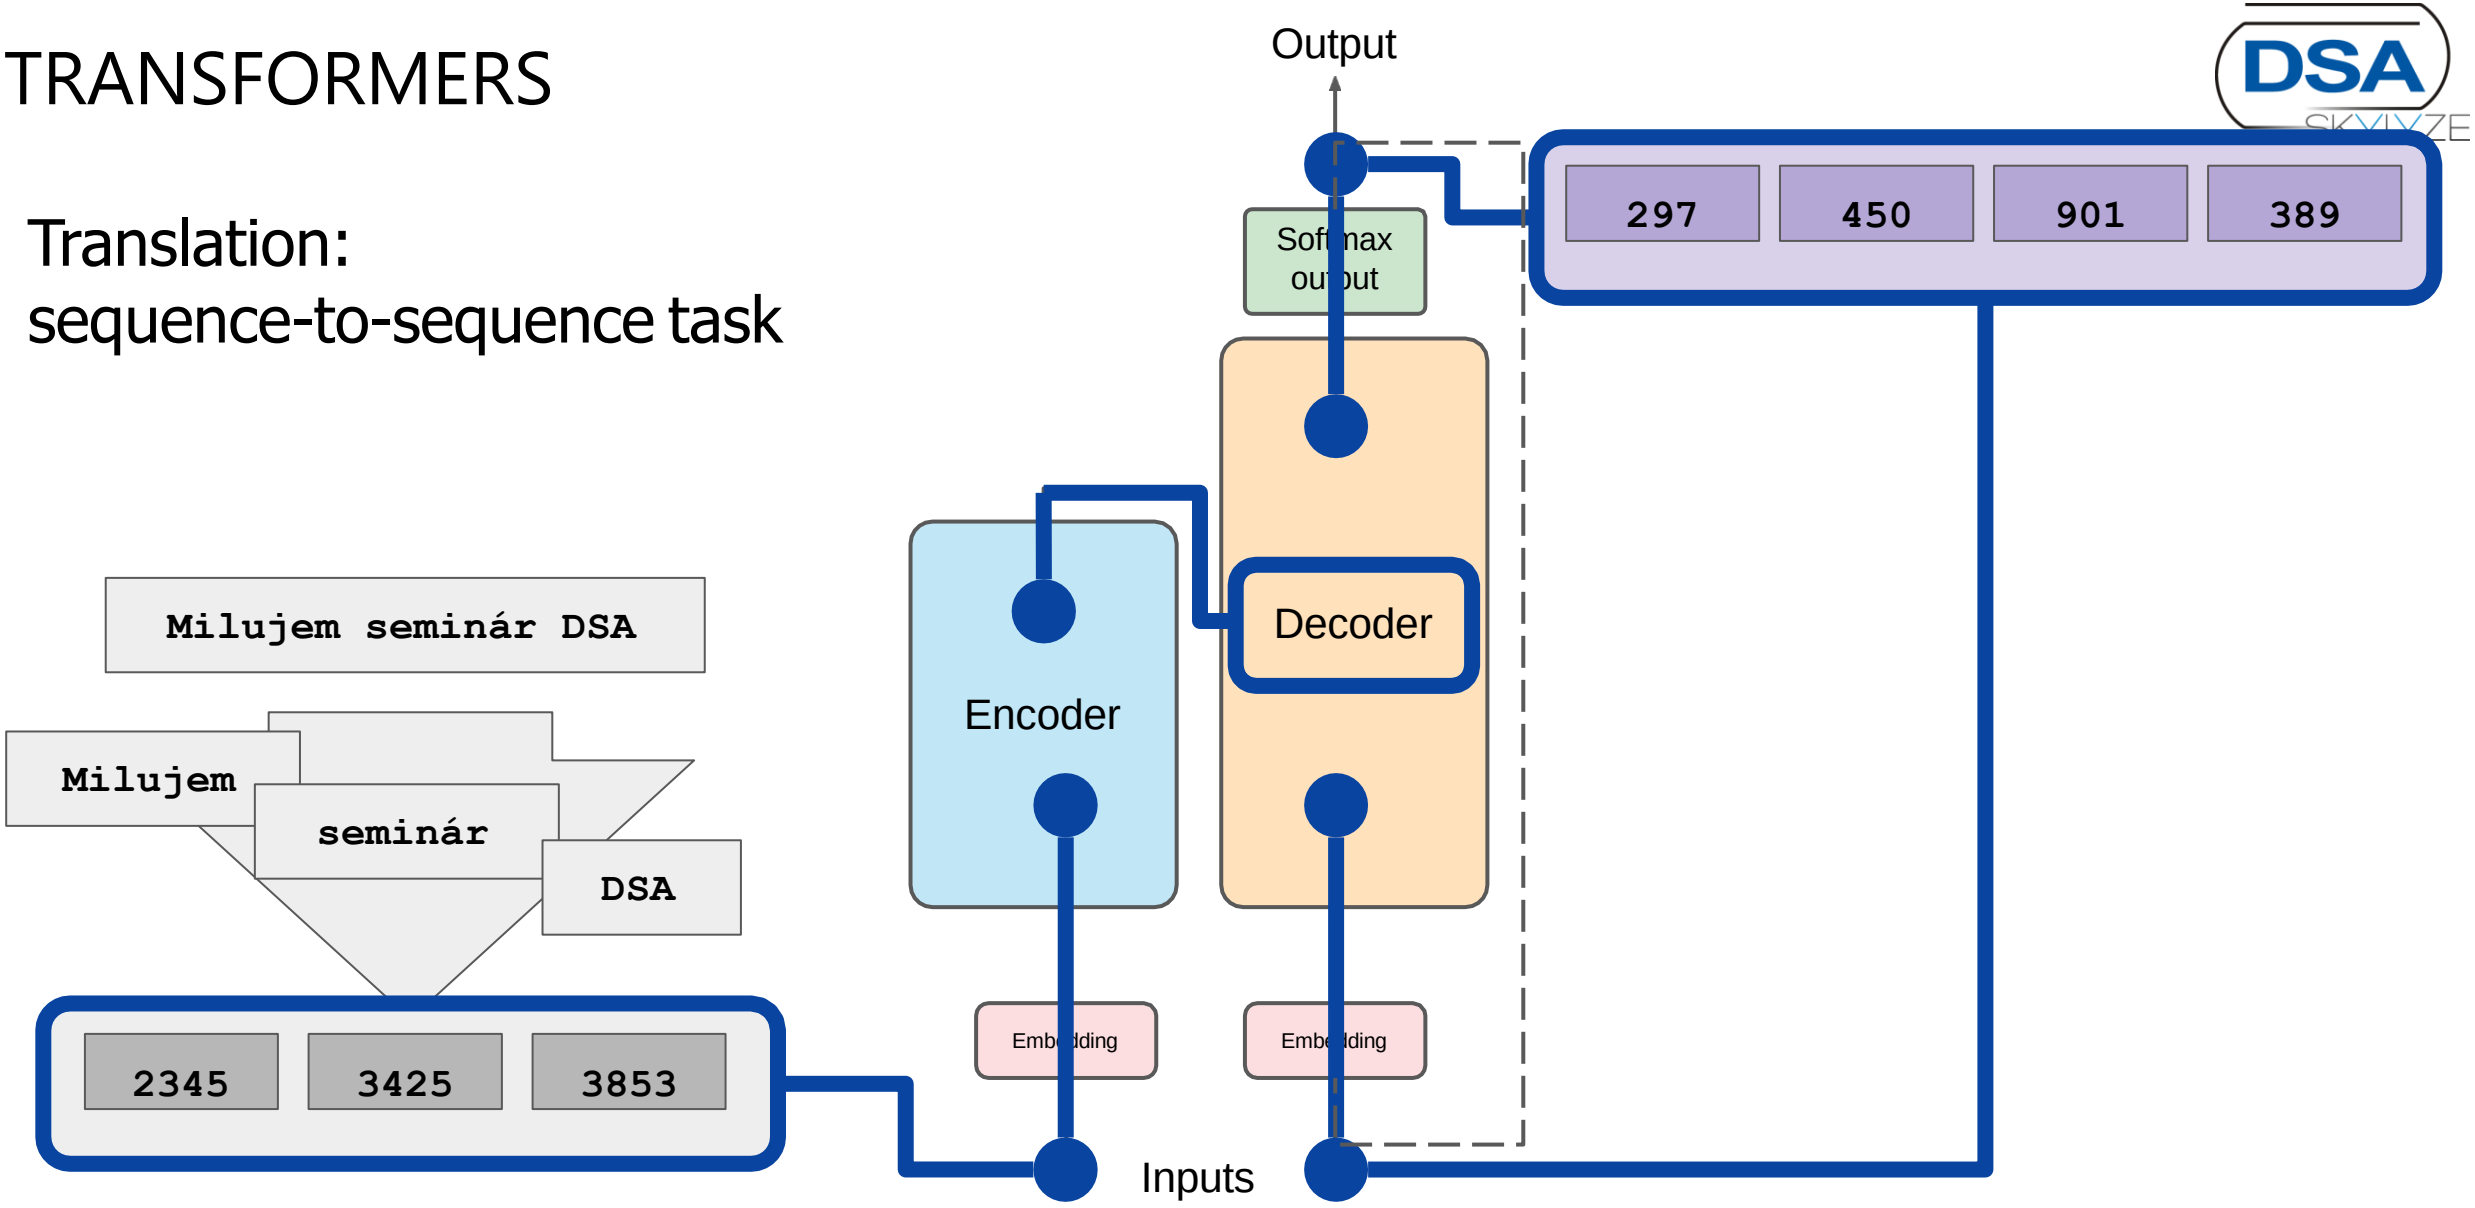

TRANSFORMERS

Translation: sequence-to-sequence task

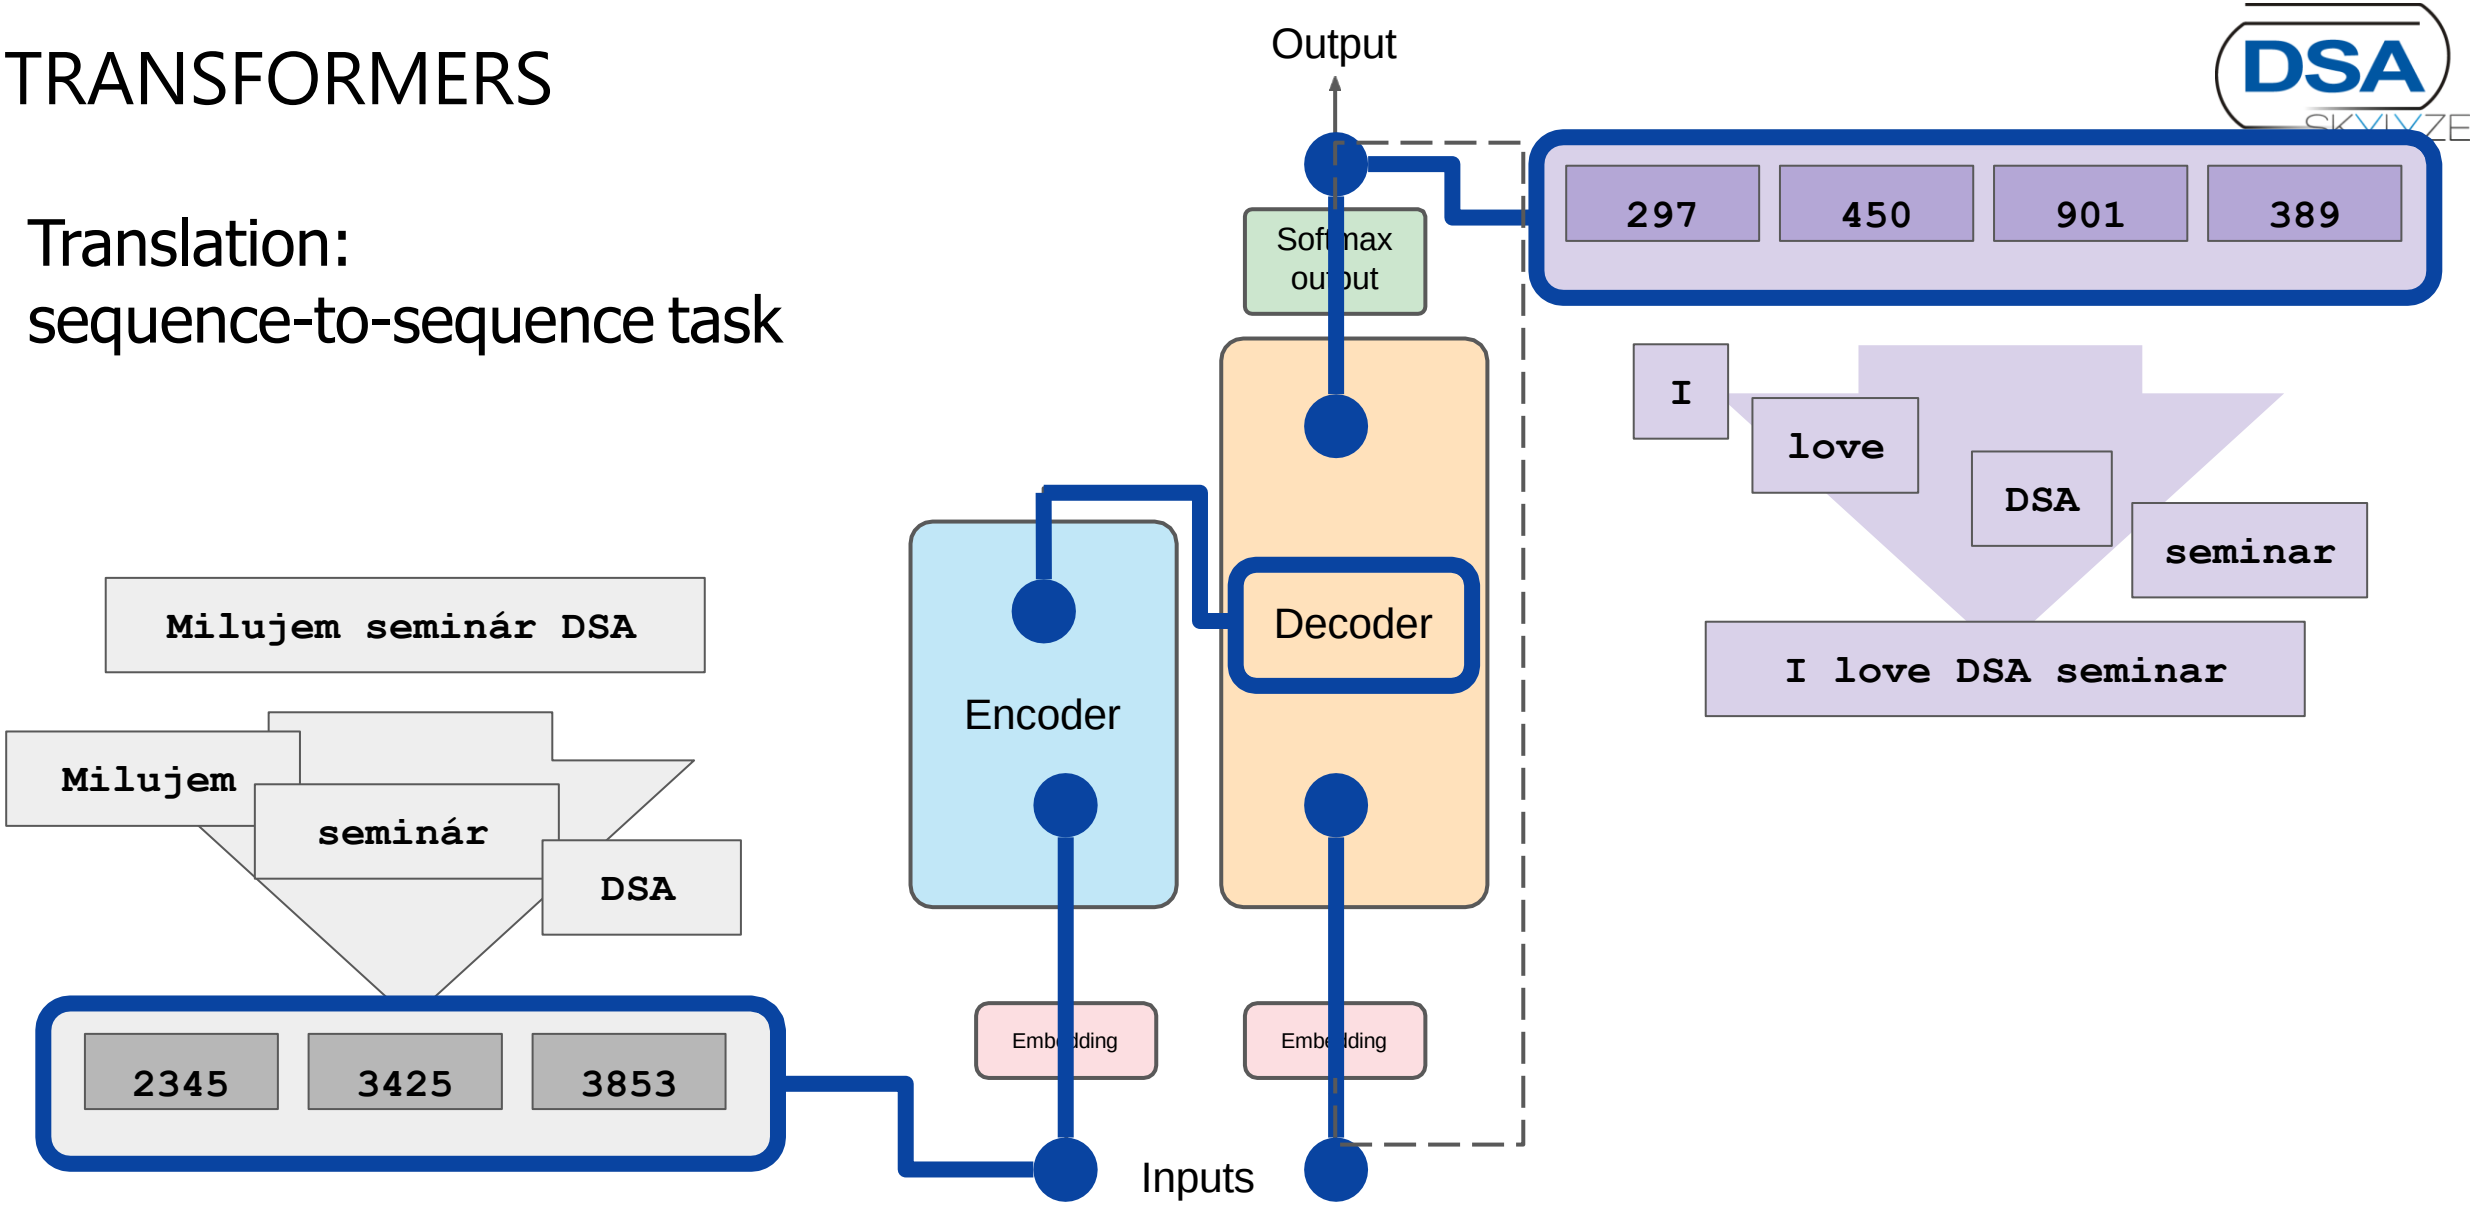

LLM USE CASES & TASKS

Essay Writing

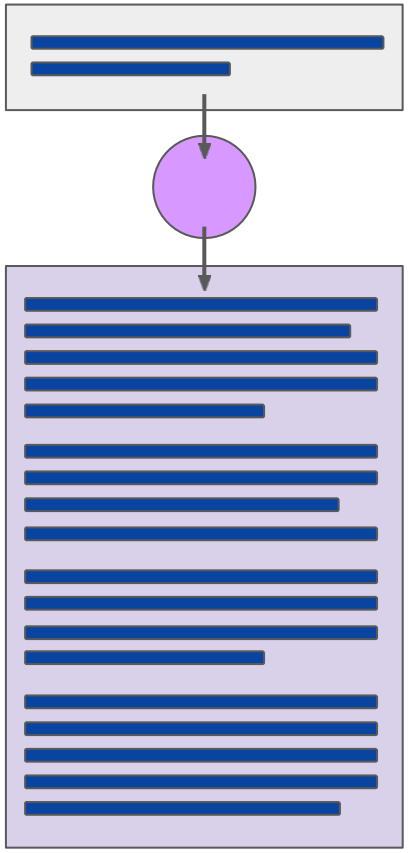

Summarization

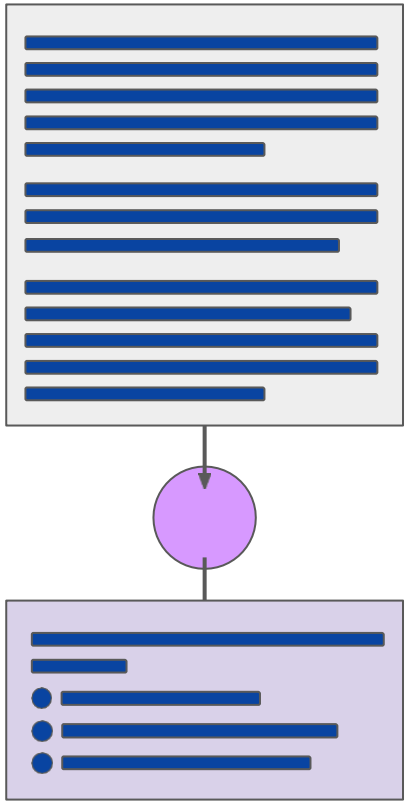

Translation

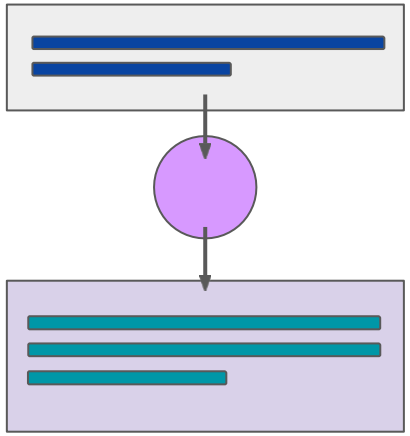

LLM USE CASES & TASKS

Essay Writing

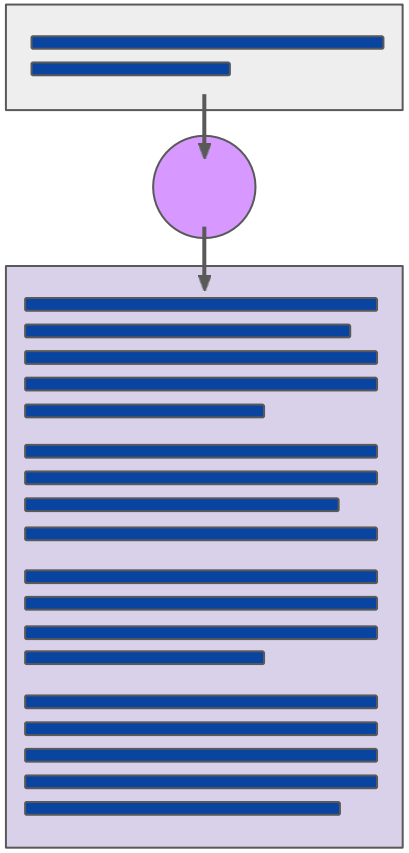

Summarization

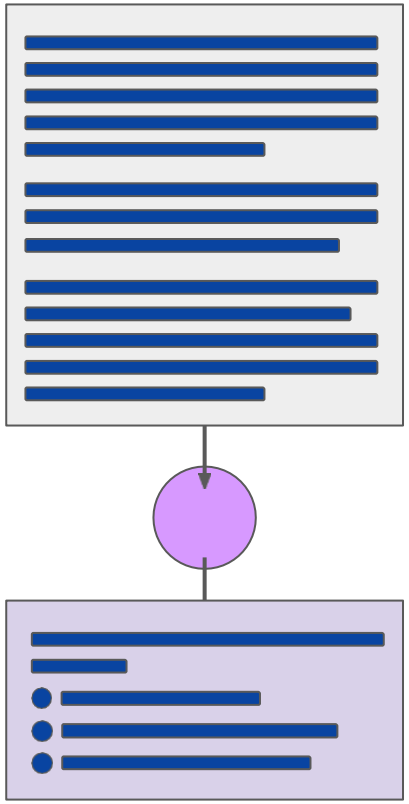

Translation

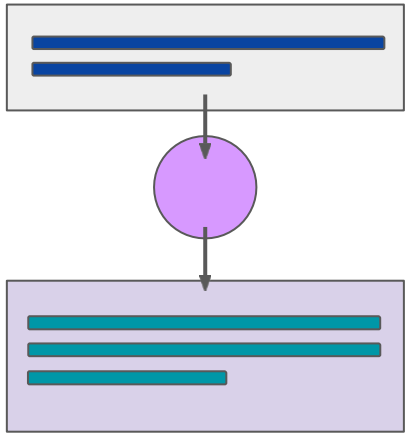

Information retrieval

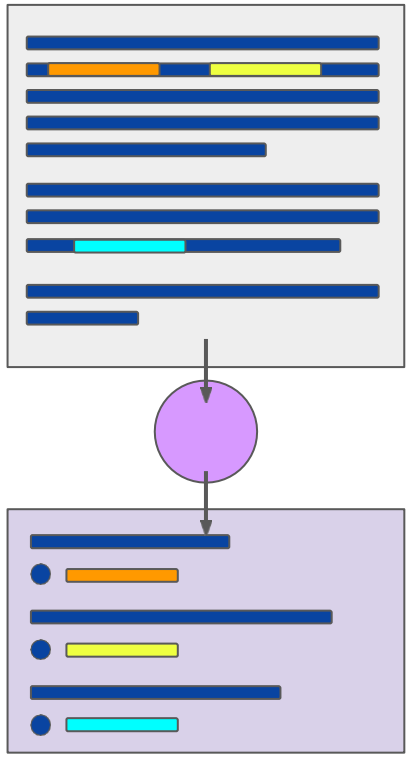

Invoke APIs and actions

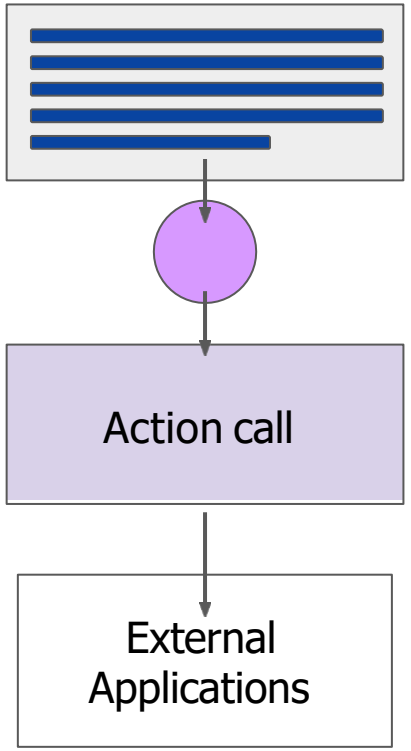

GENERATIVE AI PROJECT LIFECYCLE

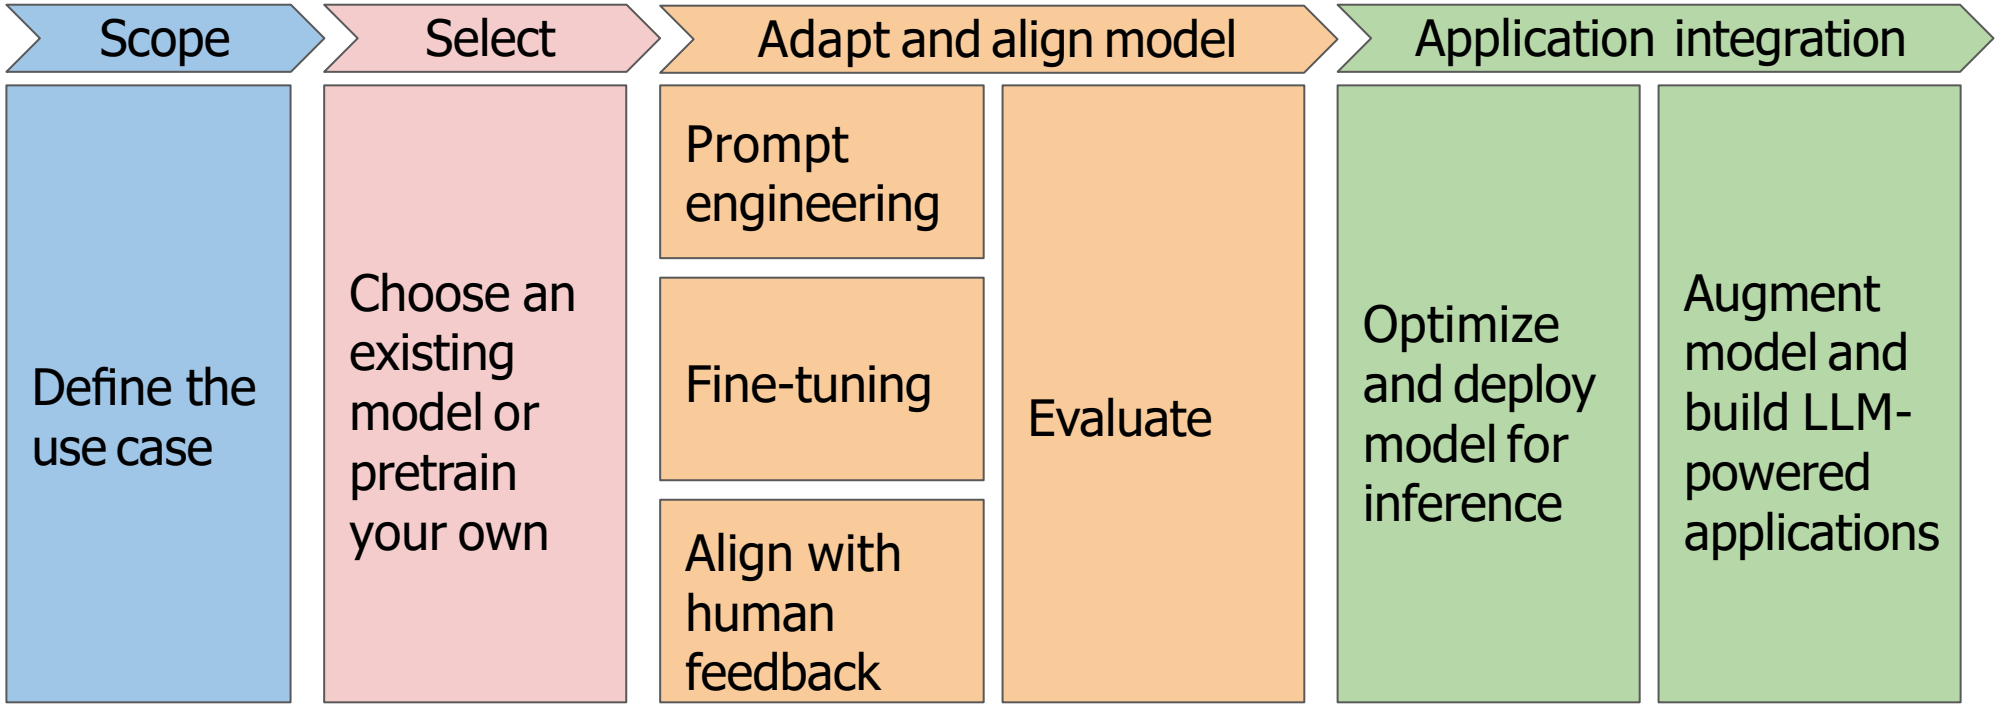

HANDS-ON

Nächstes Kapitel | Next Topic | Chapitre Prochaine | Capitolo Successivo

GENERATIVE AI PROJECT LIFECYCLE

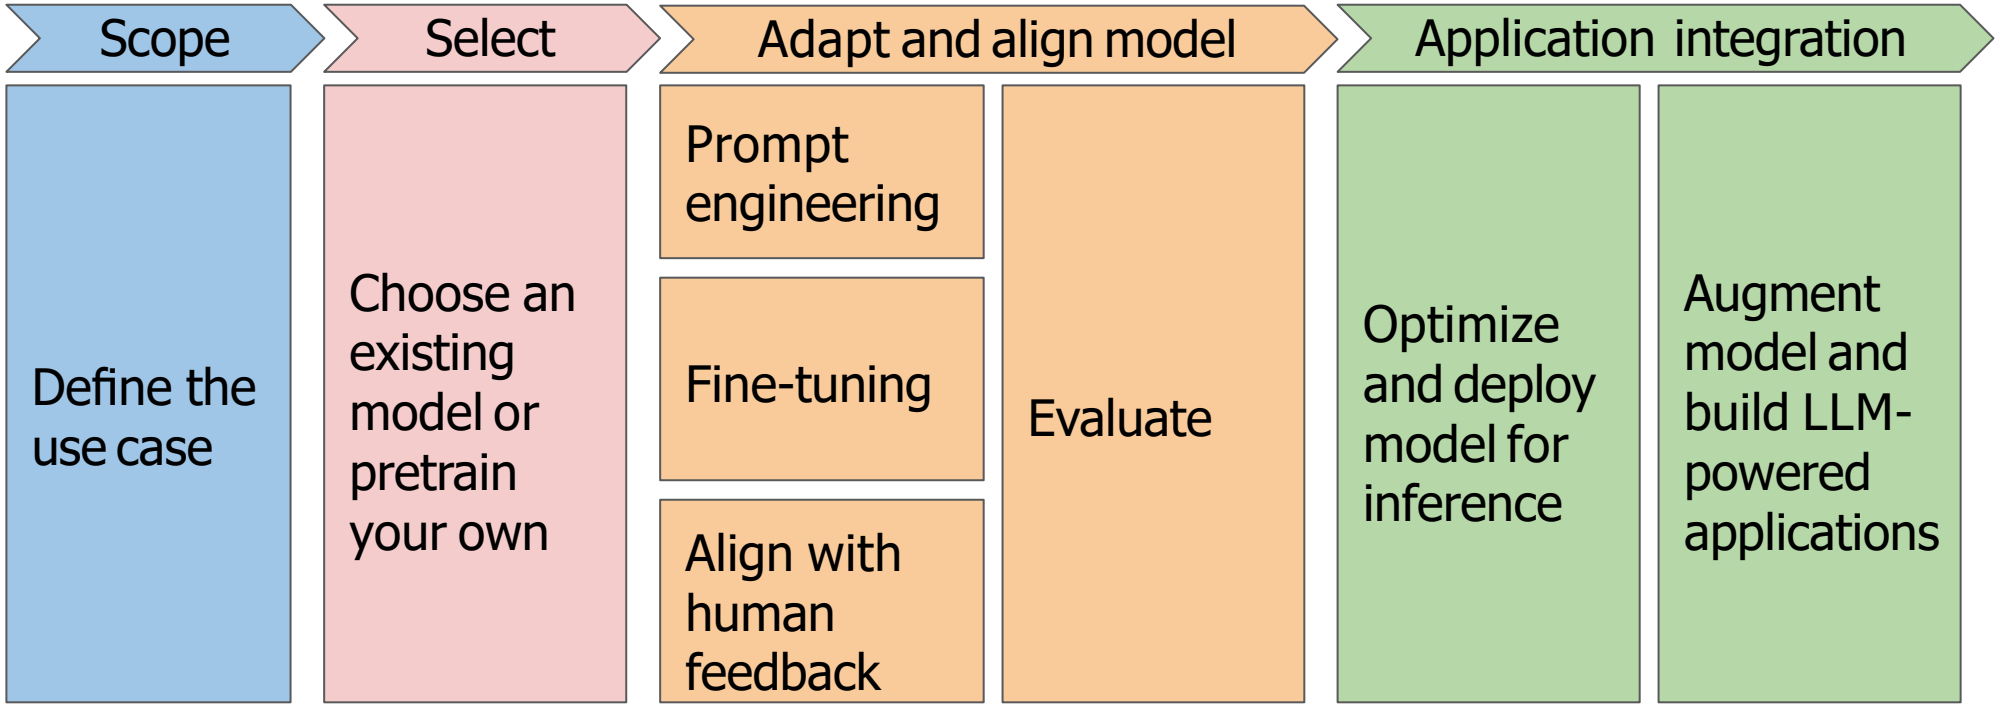

CHALLENGES OF TRAINING LLMS

Infrastructure and costs

- Trained on massive corpus of text data
- Billions to trillions of parameter
- GPU-Hardware/Infrastructure needed
- Approximate GPU RAM needed to store 1B parameters
  - 1 parameter = 4 bytes (32-bit float)
  - 1B parameters = 4 x 1e-09 bytes = 4GB
- Training needs ~20 extra bytes per parameter

APPROXIMATE GPU RAM NEEDED TO TRAIN 1B-PARAMS

DSA SKYLYZE

Memory needed to store model

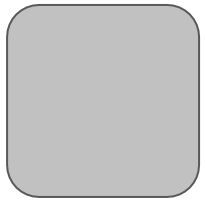

4GB @ 32-bit full precision

Memory needed to train model

GPU RAM NEEDED TO TRAIN LARGER MODELS

1B param model

175B param model
(GPT-3)

14,000 GB @ 32-bit full precision

500B param model
(PaML)

40,000 GB @ 32-bit full precision

CHALLENGES OF TRAINING LLMS

Infrastructure and costs

- It would take 288 years to train GPT-3 (175B) on a single NVIDIA Tesla V100 GPU.
  - 1Y = 8760h
  - 275W * 8760h = 2409000Wh = 2490kWh
  - 2 Person-Household : 2000 – 2800kWh
- Researchers calculated that OpenAI could have trained GPT-3 in as little as 34 days on 1,024x A100 GPUs
- It is estimated that GPT-3 cost around $4.6 million dollars to train from scratch
- Cost of this workshop ~90$ (1$/min)

- **Jun 2017**
- **16GB RAM**
- **250-300 Watt**

- **Jun 2021**
- **80GB RAM**
- **300-400 Watt**

GENERATIVE AI

Nächstes Kapitel | Next Topic | Chapitre Prochaine | Capitolo Successivo

GENERATIVE AI

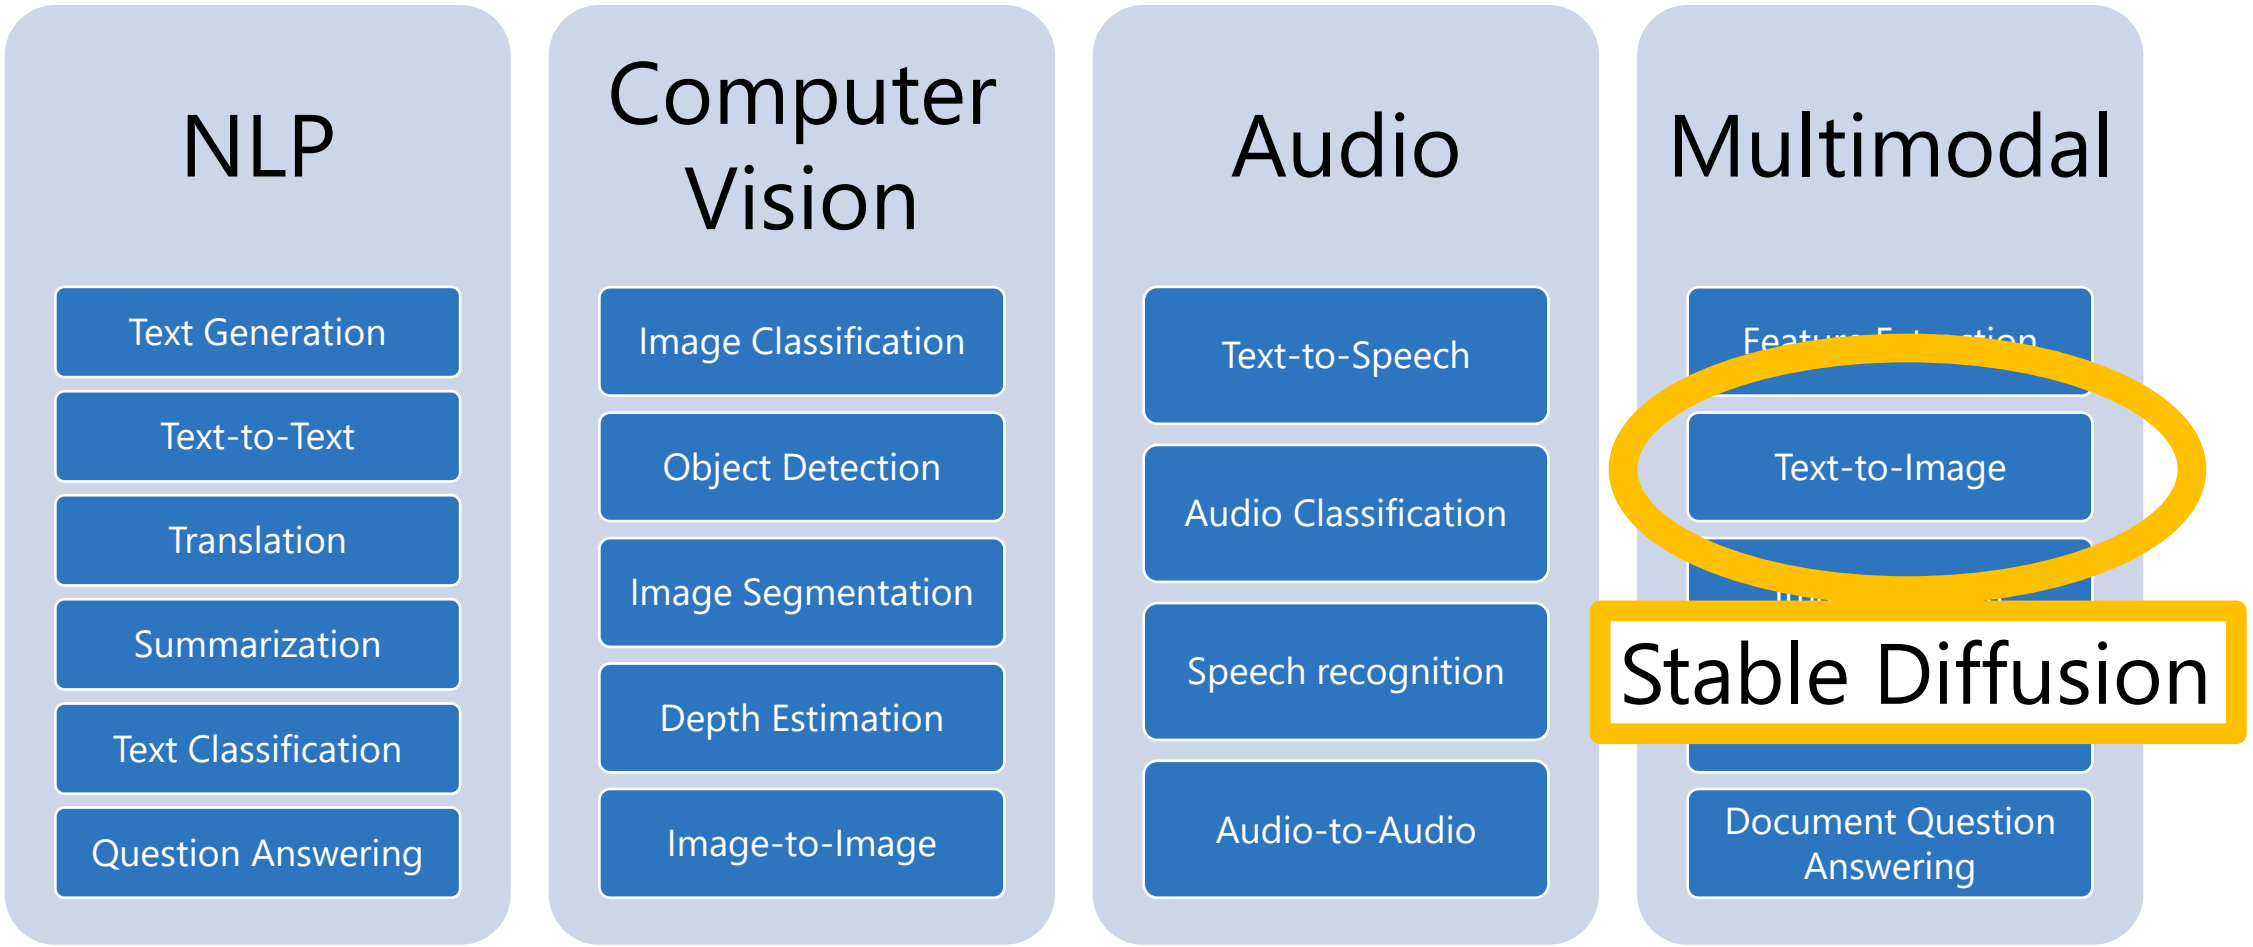

STABLE DIFFUSION

Nächstes Kapitel | Next Topic | Chapitre Prochaine | Capitolo Successivo

ACKNOWLEDGEMENT AND COPYRIGHT NOTICE

The presentation is based on the course Generative AI with Large Language Models. The slides are based on the course material from DeepLearning.AI